

Topic 3 - Further Python (3 hours)

July 5, 2022

#

Computational Statistics with Python

##

Topic 3: Further Python

##

Expected lecture time: 2-3 hours

Giancarlo Manzi

Department of Economics, Management and Quantitative Methods

University of Milan, Milan, Italy

0.1 The Pandas series object

Series is a one-dimensional labeled array from the library Pandas capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.).

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
# Using Numpy's pseudo random number generator
import numpy as np
data = np.random.randn(20)
index = range(1990, 2010)
```

```
[3]: print (data)
print (list(index))
```

```
[ 1.07412852  0.03522929 -0.00307068 -0.7969454   0.08838703 -0.09938906
 -0.56601135  0.98303324 -1.46928736  1.53688249  0.18417489  0.61477394
 -0.12605095  1.60410251  0.74209285  0.74957552 -0.312181   -0.46701963
  1.1470691  -1.22639548]
```

```
[1990, 1991, 1992, 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2000, 2001, 2002,
2003, 2004, 2005, 2006, 2007, 2008, 2009]
```

```
[4]: y = pd.Series(data, index=index)
```

```
[5]: print (y)
```

```
1990    1.074129
1991    0.035229
1992   -0.003071
1993   -0.796945
1994    0.088387
1995   -0.099389
1996   -0.566011
1997    0.983033
1998   -1.469287
1999    1.536882
2000    0.184175
2001    0.614774
2002   -0.126051
2003    1.604103
2004    0.742093
2005    0.749576
2006   -0.312181
2007   -0.467020
2008    1.147069
2009   -1.226395
dtype: float64
```

```
[6]: salaries = {
      'juan': 1500, 'maria': 2560.34, 'cesc': None, 'juan carlos': 2451
    }
```

```
[7]: s = pd.Series(salaries)
```

```
[7]: print (s)
```

```
juan          1500.00
maria          2560.34
cesc              NaN
juan carlos    2451.00
dtype: float64
```

0.1.1 Access series as arrays

```
[9]: print (s[:2])
      print (s[s > s.median()], '\n')
      print (np.log(s), '\n')
      print (s + s, '\n')
      print (s * 3, '\n')
      print (y[4:8] + y[4:10])
```

```
juan          1500.00
```

```
maria      2560.34
dtype: float64
maria      2560.34
dtype: float64
```

```
juan          7.313220
maria          7.847895
cesc           NaN
juan carlos    7.804251
dtype: float64
```

```
juan          3000.00
maria          5120.68
cesc           NaN
juan carlos    4902.00
dtype: float64
```

```
juan          4500.00
maria          7681.02
cesc           NaN
juan carlos    7353.00
dtype: float64
```

```
1994    0.176774
1995   -0.198778
1996   -1.132023
1997    1.966066
1998           NaN
1999           NaN
dtype: float64
```

0.1.2 Difference between Python list and Pandas series

```
[10]: my_list = ['a', 'b', 'c', 'd']
      print(my_list)
```

```
['a', 'b', 'c', 'd']
```

```
[11]: my_series = pd.Series(my_list, index = [2,1,3,0])
      print(my_series)
```

```
2    a
1    b
3    c
0    d
dtype: object
```

```
[14]: print(my_series[3])
```

c

```
[15]: print(my_list[3])
```

d

1 Data Frames

From <http://pandas.pydata.org/pandas-docs/stable/dsintro.html#dataframe>

DataFrame is a 2-dimensional labeled data structure with columns of potentially different types. You can think of it like a spreadsheet or SQL table, or a dict of Series objects. It is generally the most commonly used pandas object. Like Series, DataFrame accepts many different kinds of input:

- Dict of 1D ndarrays, lists, dicts, or Series
- 2-D numpy.ndarray
- Structured or record ndarray
- A series
- Another data frame

Along with the data, you can optionally pass index (row labels) and columns (column labels) arguments. If you pass an index and / or columns, you are guaranteeing the index and / or columns of the resulting DataFrame. Thus, a dict of Series plus a specific index will discard all data not matching up to the passed index.

If axis labels are not passed, they will be constructed from the input data based on common sense rules.

```
[16]: y = 2020
      s = 2000
      k = {'Smith': y, 'McDonald': s}
      k
```

```
[16]: {'Smith': 2020, 'McDonald': 2000}
```

```
[18]: df = pd.DataFrame(k.items())
      df
```

```
[18]:
```

	0	1
0	Smith	2020
1	McDonald	2000

```
[19]: print(df)
```

	0	1
0	Smith	2020
1	McDonald	2000

```
[20]: pd.DataFrame(k.items(), columns=['Name', 'Salary'])
```

```
[20]:      Name  Salary
      0    Smith    2020
      1 McDonald    2000
```

```
[21]: s = pd.Series(k, name='DateValue')
```

```
[22]: s.index.name = 'Name'
      s
```

```
[22]: Name
      Smith      2020
      McDonald  2000
      Name: DateValue, dtype: int64
```

1.1 Loading and manipulating data

Retrieve the complete local dataset from Kaggle website.

```
[27]: accidents = 'accidents_2012_to_2014.csv'
      A = pd.read_csv(accidents, low_memory=False, index_col=0)
      A
```

```
[27]:      Location_Easting_OSGR  Location_Northing_OSGR  Longitude \
Accident_Index
201201BS70001      527200      178760 -0.169101
201201BS70002      524930      181430 -0.200838
201201BS70003      525860      178080 -0.188636
201201BS70004      524980      181030 -0.200259
201201BS70005      526170      179200 -0.183773
...
2.01E+12      310037      597647 -3.417278
2.01E+12      321509      574063 -3.230255
2.01E+12      321337      566365 -3.230826
2.01E+12      323869      566853 -3.191397
2.01E+12      314072      579971 -3.348426

      Latitude  Police_Force  Accident_Severity \
Accident_Index
201201BS70001  51.493429      1      3
201201BS70002  51.517931      1      3
201201BS70003  51.487618      1      3
201201BS70004  51.514325      1      3
201201BS70005  51.497614      1      3
...
2.01E+12      55.264773      98      2
2.01E+12      55.054855      98      3
2.01E+12      54.985668      98      3
```

2.01E+12	54.990446	98	2
2.01E+12	55.106700	98	3

	Number_of_Vehicles	Number_of_Casualties	Date \
Accident_Index			
201201BS70001	2	1	19/01/2012
201201BS70002	2	1	04/01/2012
201201BS70003	2	1	10/01/2012
201201BS70004	1	1	18/01/2012
201201BS70005	1	1	17/01/2012
...
2.01E+12	2	1	07/12/2014
2.01E+12	2	2	11/12/2014
2.01E+12	1	1	09/12/2014
2.01E+12	3	2	17/12/2014
2.01E+12	2	2	24/12/2014

	Day_of_Week	...	Pedestrian_Crossing-Physical_Facilities \
Accident_Index			
201201BS70001	5	...	Pedestrian phase at traffic signal junction
201201BS70002	4	...	No physical crossing within 50 meters
201201BS70003	3	...	non-junction pedestrian crossing
201201BS70004	4	...	No physical crossing within 50 meters
201201BS70005	3	...	No physical crossing within 50 meters
...
2.01E+12	1	...	No physical crossing within 50 meters
2.01E+12	5	...	No physical crossing within 50 meters
2.01E+12	3	...	No physical crossing within 50 meters
2.01E+12	4	...	No physical crossing within 50 meters
2.01E+12	4	...	No physical crossing within 50 meters

	Light_Conditions \
Accident_Index	
201201BS70001	Darkness: Street lights present and lit
201201BS70002	Darkness: Street lights present and lit
201201BS70003	Daylight: Street light present
201201BS70004	Daylight: Street light present
201201BS70005	Darkness: Street lights present and lit
...	...
2.01E+12	Darkness: No street lighting
2.01E+12	Darkness: No street lighting
2.01E+12	Darkness: Street lights present and lit
2.01E+12	Darkness: No street lighting
2.01E+12	Daylight: Street light present

	Weather_Conditions	Road_Surface_Conditions \
Accident_Index		

201201BS70001	Fine without high winds	Dry
201201BS70002	Fine without high winds	Dry
201201BS70003	Fine without high winds	Dry
201201BS70004	Fine without high winds	Dry
201201BS70005	Fine without high winds	Dry
...
2.01E+12	Snowing without high winds	Snow
2.01E+12	Fine without high winds	Snow
2.01E+12	Fine without high winds	Frost/Ice
2.01E+12	Raining without high winds	Wet/Damp
2.01E+12	Fine without high winds	Wet/Damp

Special_Conditions_at_Site Carriageway_Hazards \

Accident_Index

201201BS70001	None	None
201201BS70002	None	None
201201BS70003	None	None
201201BS70004	None	None
201201BS70005	None	None
...
2.01E+12	None	None
2.01E+12	None	None
2.01E+12	None	None
2.01E+12	None	None
2.01E+12	None	None

Urban_or_Rural_Area \

Accident_Index

201201BS70001	1
201201BS70002	1
201201BS70003	1
201201BS70004	1
201201BS70005	1
...	...
2.01E+12	2
2.01E+12	2
2.01E+12	2
2.01E+12	2
2.01E+12	2

Did_Police_Officer_Attend_Scene_of_Accident \

Accident_Index

201201BS70001	Yes
201201BS70002	Yes
201201BS70003	Yes
201201BS70004	Yes
201201BS70005	Yes

```

...
2.01E+12
2.01E+12
2.01E+12
2.01E+12
2.01E+12
2.01E+12

```

```

LSOA_of_Accident_Location Year
Accident_Index
201201BS70001 E01002821 2012
201201BS70002 E01004760 2012
201201BS70003 E01002893 2012
201201BS70004 E01002886 2012
201201BS70005 E01002890 2012
...
2.01E+12 NaN 2014
2.01E+12 NaN 2014
2.01E+12 NaN 2014
2.01E+12 NaN 2014
2.01E+12 NaN 2014

```

[464697 rows x 32 columns]

[28]: A.head()

```

[28]: Location_Easting_OSGR Location_Northing_OSGR Longitude \
Accident_Index
201201BS70001 527200 178760 -0.169101
201201BS70002 524930 181430 -0.200838
201201BS70003 525860 178080 -0.188636
201201BS70004 524980 181030 -0.200259
201201BS70005 526170 179200 -0.183773

```

```

Latitude Police_Force Accident_Severity \
Accident_Index
201201BS70001 51.493429 1 3
201201BS70002 51.517931 1 3
201201BS70003 51.487618 1 3
201201BS70004 51.514325 1 3
201201BS70005 51.497614 1 3

```

```

Number_of_Vehicles Number_of_Casualties Date \
Accident_Index
201201BS70001 2 1 19/01/2012
201201BS70002 2 1 04/01/2012
201201BS70003 2 1 10/01/2012
201201BS70004 1 1 18/01/2012

```


201201BS70005

1

1 17/01/2012

Accident_Index	Day_of_Week	...	Pedestrian_Crossing-Physical_Facilities	\
201201BS70001	5	...	Pedestrian phase at traffic signal junction	
201201BS70002	4	...	No physical crossing within 50 meters	
201201BS70003	3	...	non-junction pedestrian crossing	
201201BS70004	4	...	No physical crossing within 50 meters	
201201BS70005	3	...	No physical crossing within 50 meters	

Accident_Index	Light_Conditions	\
201201BS70001	Darkness: Street lights present and lit	
201201BS70002	Darkness: Street lights present and lit	
201201BS70003	Daylight: Street light present	
201201BS70004	Daylight: Street light present	
201201BS70005	Darkness: Street lights present and lit	

Accident_Index	Weather_Conditions	Road_Surface_Conditions	\
201201BS70001	Fine without high winds	Dry	
201201BS70002	Fine without high winds	Dry	
201201BS70003	Fine without high winds	Dry	
201201BS70004	Fine without high winds	Dry	
201201BS70005	Fine without high winds	Dry	

Accident_Index	Special_Conditions_at_Site	Carriageway_Hazards	\
201201BS70001	None	None	
201201BS70002	None	None	
201201BS70003	None	None	
201201BS70004	None	None	
201201BS70005	None	None	

Accident_Index	Urban_or_Rural_Area	\
201201BS70001	1	
201201BS70002	1	
201201BS70003	1	
201201BS70004	1	
201201BS70005	1	

Accident_Index	Did_Police_Officer_Attend_Scene_of_Accident	\
201201BS70001	Yes	
201201BS70002	Yes	
201201BS70003	Yes	

201201BS70004	Yes
201201BS70005	Yes

	LSOA_of_Accident_Location	Year
Accident_Index		
201201BS70001	E01002821	2012
201201BS70002	E01004760	2012
201201BS70003	E01002893	2012
201201BS70004	E01002886	2012
201201BS70005	E01002890	2012

[5 rows x 32 columns]

```
[29]: A[['Date', 'Time']].head()
```

```
[29]:
```

	Date	Time
Accident_Index		
201201BS70001	19/01/2012	20:35
201201BS70002	04/01/2012	17:00
201201BS70003	10/01/2012	10:07
201201BS70004	18/01/2012	12:20
201201BS70005	17/01/2012	20:24

```
[23]: A.dtypes
```

```
[23]:
```

Location_Easting_OSGR	int64
Location_Northing_OSGR	int64
Longitude	float64
Latitude	float64
Police_Force	int64
Accident_Severity	int64
Number_of_Vehicles	int64
Number_of_Casualties	int64
Date	object
Day_of_Week	int64
Time	object
Local_Authority_(District)	int64
Local_Authority_(Highway)	object
1st_Road_Class	int64
1st_Road_Number	int64
Road_Type	object
Speed_limit	int64
Junction_Detail	float64
Junction_Control	object
2nd_Road_Class	int64
2nd_Road_Number	int64
Pedestrian_Crossing-Human_Control	object

Pedestrian_Crossing-Physical_Facilities	object
Light_Conditions	object
Weather_Conditions	object
Road_Surface_Conditions	object
Special_Conditions_at_Site	object
Carriageway_Hazards	object
Urban_or_Rural_Area	int64
Did_Police_Officer_Attend_Scene_of_Accident	object
LSOA_of_Accident_Location	object
Year	int64
dtype:	object

```
[30]: from datetime import datetime

def todate(d, t):
    try:
        dt = datetime.strptime(" ".join([d, t]), '%d/%m/%Y %H:%M')
    except TypeError:
        dt = np.nan
    return dt
```

```
[31]: A['Datetime'] = [todate(x.Date, x.Time) for i, x in A.iterrows()]
```

```
[32]: A[['Datetime', 'Police_Force']].head()
```

```
[32]:
```

	Datetime	Police_Force
Accident_Index		
201201BS70001	2012-01-19 20:35:00	1
201201BS70002	2012-01-04 17:00:00	1
201201BS70003	2012-01-10 10:07:00	1
201201BS70004	2012-01-18 12:20:00	1
201201BS70005	2012-01-17 20:24:00	1

```
[33]: A.shape
```

```
[33]: (464697, 33)
```

```
[34]: A.dtypes
```

```
[34]:
```

Location_Easting_OSGR	int64
Location_Northing_OSGR	int64
Longitude	float64
Latitude	float64
Police_Force	int64
Accident_Severity	int64
Number_of_Vehicles	int64
Number_of_Casualties	int64

Date	object
Day_of_Week	int64
Time	object
Local_Authority_(District)	int64
Local_Authority_(Highway)	object
1st_Road_Class	int64
1st_Road_Number	int64
Road_Type	object
Speed_limit	int64
Junction_Detail	float64
Junction_Control	object
2nd_Road_Class	int64
2nd_Road_Number	int64
Pedestrian_Crossing-Human_Control	object
Pedestrian_Crossing-Physical_Facilities	object
Light_Conditions	object
Weather_Conditions	object
Road_Surface_Conditions	object
Special_Conditions_at_Site	object
Carriageway_Hazards	object
Urban_or_Rural_Area	int64
Did_Police_Officer_Attend_Scene_of_Accident	object
LSOA_of_Accident_Location	object
Year	int64
Datetime	datetime64[ns]
dtype:	object

1.2 Access dataframe by index and col

```
[35]: my_df = A.iloc[2:6] # gets rows (or columns) at particular positions in the
      ↪ index (so it only takes integers).
      my_df
```

```
[35]:
```

	Location_Easting_OSGR	Location_Northing_OSGR	Longitude	\
Accident_Index				
201201BS70003	525860	178080	-0.188636	
201201BS70004	524980	181030	-0.200259	
201201BS70005	526170	179200	-0.183773	
201201BS70006	526090	177600	-0.185496	

	Latitude	Police_Force	Accident_Severity	\
Accident_Index				
201201BS70003	51.487618	1	3	
201201BS70004	51.514325	1	3	
201201BS70005	51.497614	1	3	
201201BS70006	51.483253	1	3	

Accident_Index	Number_of_Vehicles	Number_of_Casualties	Date \
201201BS70003	2	1	10/01/2012
201201BS70004	1	1	18/01/2012
201201BS70005	1	1	17/01/2012
201201BS70006	2	1	19/01/2012

Accident_Index	Day_of_Week ...	Light_Conditions \
201201BS70003	3 ...	Daylight: Street light present
201201BS70004	4 ...	Daylight: Street light present
201201BS70005	3 ...	Darkness: Street lights present and lit
201201BS70006	5 ...	Darkness: Street lights present and lit

Accident_Index	Weather_Conditions	Road_Surface_Conditions \
201201BS70003	Fine without high winds	Dry
201201BS70004	Fine without high winds	Dry
201201BS70005	Fine without high winds	Dry
201201BS70006	Raining without high winds	Wet/Damp

Accident_Index	Special_Conditions_at_Site	Carriageway_Hazards \
201201BS70003	None	None
201201BS70004	None	None
201201BS70005	None	None
201201BS70006	None	None

Accident_Index	Urban_or_Rural_Area \
201201BS70003	1
201201BS70004	1
201201BS70005	1
201201BS70006	1

Accident_Index	Did_Police_Officer_Attend_Scene_of_Accident \
201201BS70003	Yes
201201BS70004	Yes
201201BS70005	Yes
201201BS70006	Yes

Accident_Index	LSOA_of_Accident_Location	Year	Datetime
201201BS70003	E01002893	2012	2012-01-10 10:07:00
201201BS70004	E01002886	2012	2012-01-18 12:20:00
201201BS70005	E01002890	2012	2012-01-17 20:24:00

201201BS70006

E01002912 2012 2012-01-19 07:30:00

[4 rows x 33 columns]

```
[36]: #SUBSETTING a data frame
selection = A[A['Road_Surface_Conditions'] == 'Dry'].sort_values(
    'Number_of_Casualties', ascending=False)
selection
#selection[['Weather_Conditions', 'Police_Force',
#           'Accident_Severity', 'Number_of_Vehicles', 'Number_of_Casualties']].
↳head()
```

```
[36]:      Location_Easting_OSGR  Location_Northing_OSGR  Longitude \
Accident_Index
20144100J0489          523000          199780 -0.222211
201411NH11644          418196          552132 -1.718034
2.01E+12          591380          169440  0.749417
2.01E+12          375840          203065 -2.351207
201422E404170          355950          235980 -2.643371
...
201297QC00409          304120          637780 -3.524192
201297QC00510          281810          652360 -3.884585
201297QC00605          294640          612550 -3.665077
201297QC00606          302140          641540 -3.556962
2.01E+12          311812          580747 -3.384080
```

```
      Latitude  Police_Force  Accident_Severity \
Accident_Index
20144100J0489  51.683269          41          2
201411NH11644  54.863663          11          2
2.01E+12      51.391660          46          2
2.01E+12      51.725734          53          2
201422E404170  52.020443          22          2
...
201297QC00409  55.624152          97          3
201297QC00510  55.750175          97          3
201297QC00605  55.395579          97          3
201297QC00606  55.657530          97          1
2.01E+12      55.113274          98          3
```

```
      Number_of_Vehicles  Number_of_Casualties      Date \
Accident_Index
20144100J0489          2          93  20/10/2014
201411NH11644          2          87  03/06/2014
2.01E+12          67          70  05/09/2013
2.01E+12          1          54  19/08/2014
201422E404170          2          41  10/11/2014
```

...
201297QC00409	5	1	07/09/2012
201297QC00510	2	1	02/10/2012
201297QC00605	1	1	05/05/2012
201297QC00606	2	1	05/06/2012
2.01E+12	1	1	17/11/2014

	Day_of_Week	...	Light_Conditions	\
Accident_Index		...		
20144100J0489	2	...	Daylight: Street light present	
201411NH11644	3	...	Daylight: Street light present	
2.01E+12	5	...	Daylight: Street light present	
2.01E+12	3	...	Daylight: Street light present	
201422E404170	2	...	Daylight: Street light present	
...	
201297QC00409	6	...	Daylight: Street light present	
201297QC00510	3	...	Daylight: Street light present	
201297QC00605	7	...	Daylight: Street light present	
201297QC00606	3	...	Daylight: Street light present	
2.01E+12	2	...	Daylight: Street light present	

	Weather_Conditions	Road_Surface_Conditions	\
Accident_Index			
20144100J0489	Fine without high winds	Dry	
201411NH11644	Fine without high winds	Dry	
2.01E+12	Fog or mist	Dry	
2.01E+12	Fine without high winds	Dry	
201422E404170	Fine without high winds	Dry	
...	
201297QC00409	Fine without high winds	Dry	
201297QC00510	Fine without high winds	Dry	
201297QC00605	Fine without high winds	Dry	
201297QC00606	Fine without high winds	Dry	
2.01E+12	Fine without high winds	Dry	

	Special_Conditions_at_Site	Carriageway_Hazards	\
Accident_Index			
20144100J0489	None	None	
201411NH11644	None	None	
2.01E+12	None	None	
2.01E+12	None	None	
201422E404170	None	None	
...	
201297QC00409	None	None	
201297QC00510	None	None	
201297QC00605	Road surface defective	None	
201297QC00606	None	None	

2.01E+12	None	None
----------	------	------

Urban_or_Rural_Area \	Accident_Index
20144100J0489	2
201411NH11644	2
2.01E+12	2
2.01E+12	2
201422E404170	2
...	...
201297QC00409	2
201297QC00510	2
201297QC00605	2
201297QC00606	2
2.01E+12	2

Did_Police_Officer_Attend_Scene_of_Accident \	Accident_Index
20144100J0489	Yes
201411NH11644	Yes
2.01E+12	Yes
2.01E+12	Yes
201422E404170	Yes
...	...
201297QC00409	Yes
201297QC00510	Yes
201297QC00605	No
201297QC00606	Yes
2.01E+12	Yes

LSOA_of_Accident_Location	Year	Datetime
Accident_Index		
20144100J0489	E01023584 2014	2014-10-20 08:22:00
201411NH11644	E01020624 2014	2014-06-03 08:22:00
2.01E+12	E01024597 2013	2013-09-05 07:15:00
2.01E+12	E01022371 2014	2014-08-19 09:57:00
201422E404170	E01014026 2014	2014-11-10 08:20:00
...
201297QC00409	NaN 2012	2012-09-07 15:45:00
201297QC00510	NaN 2012	2012-10-02 17:45:00
201297QC00605	NaN 2012	2012-05-05 16:00:00
201297QC00606	NaN 2012	2012-06-05 15:39:00
2.01E+12	NaN 2014	2014-11-17 13:10:00

[319370 rows x 33 columns]


```
[37]: selection[['Weather_Conditions', 'Police_Force', 'Accident_Severity',
               'Number_of_Vehicles', 'Number_of_Casualties']].
       ↳groupby('Weather_Conditions').mean()
```

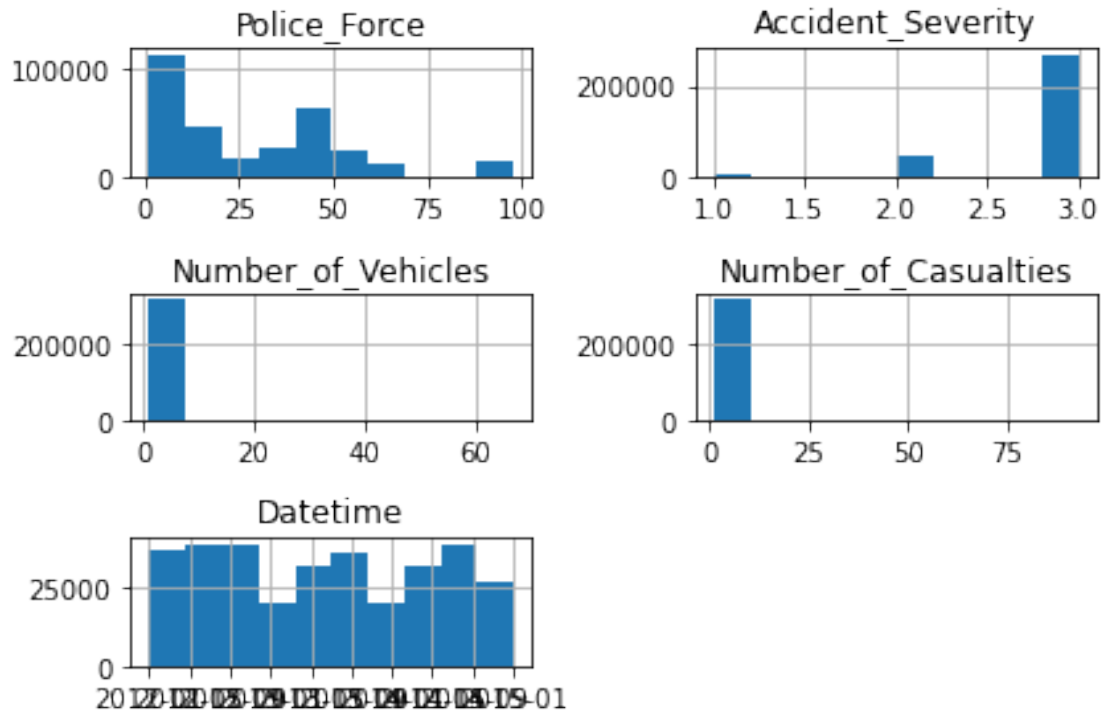
```
[37]:
```

	Police_Force	Accident_Severity \
Weather_Conditions		
Fine with high winds	32.652875	2.811360
Fine without high winds	27.051892	2.830949
Fog or mist	39.051163	2.797674
Other	29.449333	2.868000
Raining with high winds	32.687500	2.833333
Raining without high winds	38.734211	2.873684
Snowing with high winds	45.666667	2.666667
Snowing without high winds	31.560976	2.902439
Unknown	27.058422	2.872004

	Number_of_Vehicles	Number_of_Casualties
Weather_Conditions		
Fine with high winds	1.796283	1.352384
Fine without high winds	1.846165	1.321455
Fog or mist	1.997674	1.520930
Other	1.788000	1.269333
Raining with high winds	1.895833	1.458333
Raining without high winds	1.792105	1.297368
Snowing with high winds	1.777778	1.777778
Snowing without high winds	1.780488	1.195122
Unknown	1.766977	1.217710

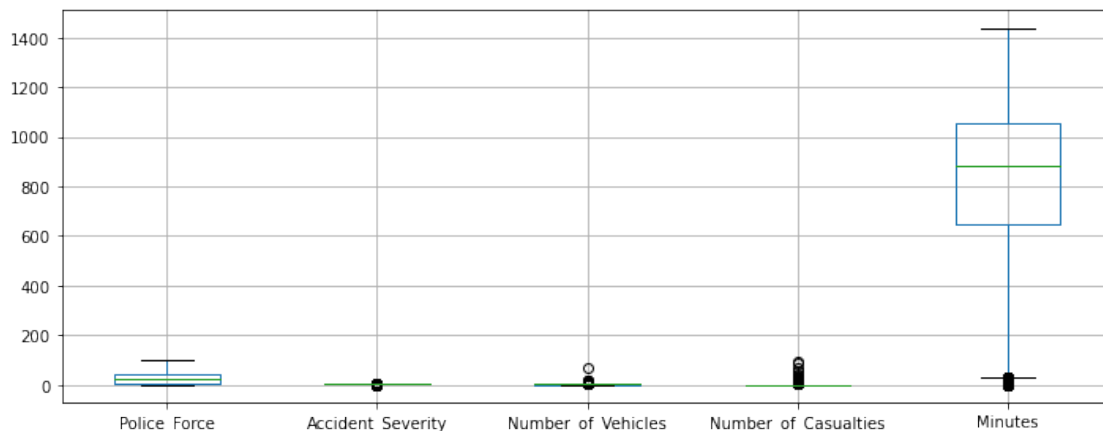
```
[38]: sel = selection[['Weather_Conditions', 'Police_Force', 'Accident_Severity',
                       'Number_of_Vehicles', 'Number_of_Casualties', 'Datetime']]
```

```
[34]: sel.hist()
      plt.tight_layout()
      plt.show()
```

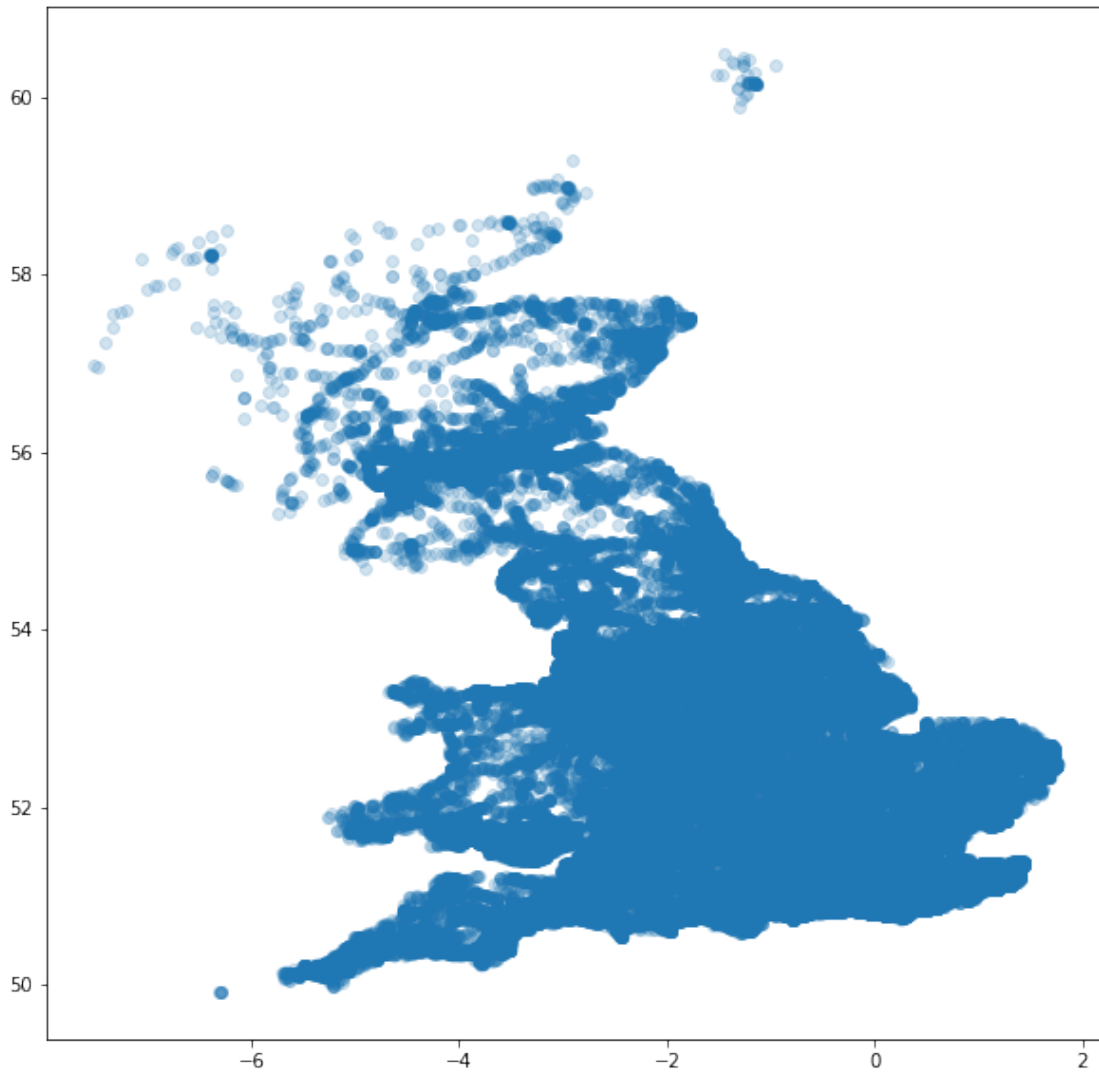


```
[40]: minutes = []
      for i, row in sel.iterrows():
          h, m = row['Datetime'].hour, row['Datetime'].minute
          minutes.append(h*60 + m)
      sel = sel.copy()
      sel['Minutes'] = minutes
```

```
[42]: fig, axes = plt.subplots(nrows=1, ncols=1, figsize=(10, 4), sharey=True)
      sel.boxplot(ax=axes)
      plt.tight_layout()
      plt.show()
```



```
[38]: fig, axes = plt.subplots(nrows=1, ncols=1, figsize=(10, 10), sharey=True)
      axes.scatter(selection.Longitude.values, selection.Latitude.values, alpha=0.2)
      plt.show()
```



```
[43]: pip install geopandas
```

Requirement already satisfied: geopandas in /opt/anaconda3/lib/python3.8/site-packages (0.9.0)
Requirement already satisfied: fiona>=1.8 in /opt/anaconda3/lib/python3.8/site-packages (from geopandas) (1.8.19)
Requirement already satisfied: shapely>=1.6 in /opt/anaconda3/lib/python3.8/site-packages (from geopandas) (1.7.1)

Requirement already satisfied: pandas>=0.24.0 in /opt/anaconda3/lib/python3.8/site-packages (from geopandas) (1.3.4)

Requirement already satisfied: pyproj>=2.2.0 in /opt/anaconda3/lib/python3.8/site-packages (from geopandas) (3.0.1)

Requirement already satisfied: attrs>=17 in /opt/anaconda3/lib/python3.8/site-packages (from fiona>=1.8->geopandas) (20.3.0)

Requirement already satisfied: cligj>=0.5 in /opt/anaconda3/lib/python3.8/site-packages (from fiona>=1.8->geopandas) (0.7.1)

Requirement already satisfied: six>=1.7 in /opt/anaconda3/lib/python3.8/site-packages (from fiona>=1.8->geopandas) (1.15.0)

Requirement already satisfied: certifi in /opt/anaconda3/lib/python3.8/site-packages (from fiona>=1.8->geopandas) (2020.12.5)

Requirement already satisfied: click-plugins>=1.0 in /opt/anaconda3/lib/python3.8/site-packages (from fiona>=1.8->geopandas) (1.1.1)

Requirement already satisfied: click<8,>=4.0 in /opt/anaconda3/lib/python3.8/site-packages (from fiona>=1.8->geopandas) (7.1.2)

Requirement already satisfied: munch in /opt/anaconda3/lib/python3.8/site-packages (from fiona>=1.8->geopandas) (2.5.0)

Requirement already satisfied: python-dateutil>=2.7.3 in /opt/anaconda3/lib/python3.8/site-packages (from pandas>=0.24.0->geopandas) (2.8.1)

Requirement already satisfied: pytz>=2017.3 in /opt/anaconda3/lib/python3.8/site-packages (from pandas>=0.24.0->geopandas) (2021.1)

Requirement already satisfied: numpy>=1.17.3 in /opt/anaconda3/lib/python3.8/site-packages (from pandas>=0.24.0->geopandas) (1.22.1)

WARNING: You are using pip version 21.2.2; however, version 22.1.2 is available.

You should consider upgrading via the '/opt/anaconda3/bin/python -m pip install --upgrade pip' command.

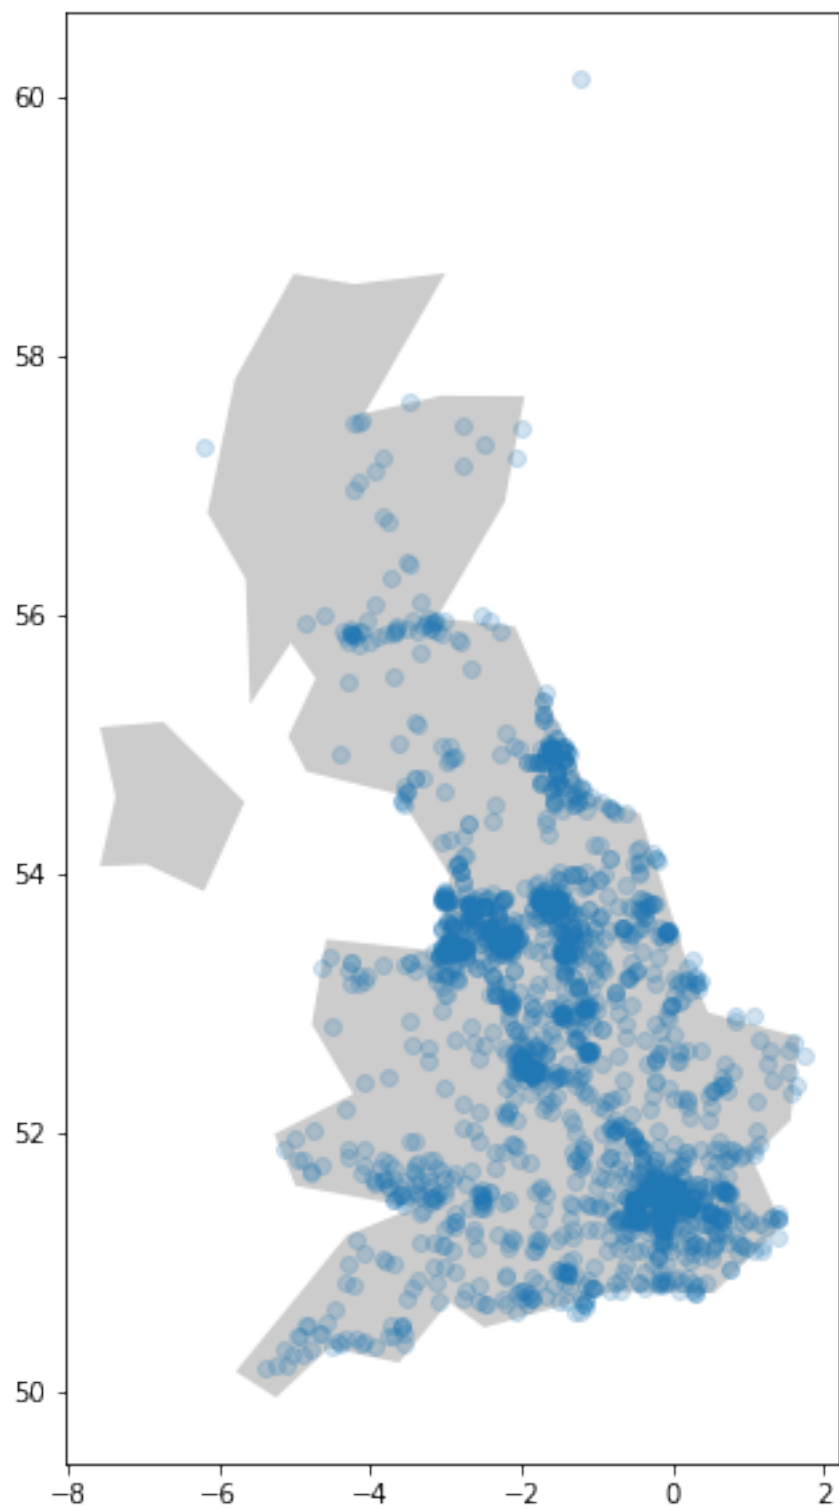
Note: you may need to restart the kernel to use updated packages.

```
[44]: import geopandas as gpd
```

```
[45]: world = gpd.read_file(gpd.datasets.get_path('naturalearth_lowres'))
```

```
[46]: UK = world[world['iso_a3']=='GBR']
```

```
[47]: limit = 2000
fig, axes = plt.subplots(nrows=1, ncols=1, figsize=(10, 10), sharey=True)
UK.plot(ax=axes, color='#CCCCCC')
axes.scatter(selection.Longitude.values[:limit], selection.Latitude.values[:
    ↪limit], alpha=0.2)
plt.show()
```



2 Standard tools for machine learning

```
[51]: # Standard import for ML
import numpy as np
import os
import tarfile
import requests
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
%matplotlib inline

# Matplotlib default setting
# When using the 'inline' backend, your matplotlib graphs will be included in_
↳ your notebook, next to the code
%matplotlib inline

mpl.rc('axes', labelsizes=14)
mpl.rc('xtick', labelsizes=12)
mpl.rc('ytick', labelsizes=12)
#np1. + any method or function you want to use
```

3 Get the data

Get the housing data (<https://www.kaggle.com/harrywang/housing>) from the Web through requests and load into a DataFrame from file

```
[ ]:
```

```
[53]: url_data = "https://raw.githubusercontent.com/ageron/handson-ml/master/datasets/
↳ housing/housing.tgz"
data_path = os.path.join("datasets", "housing")
if not os.path.isdir(data_path):
    os.makedirs(data_path)
with open(os.path.join(data_path, 'housing.tgz'), 'wb') as f:
    f.write(requests.get(url_data).content)
housing_tgz = tarfile.open(os.path.join(data_path, 'housing.tgz'))
housing_tgz.extractall(path=data_path)
housing_tgz.close()
```

```
[54]: housing = pd.read_csv('datasets/housing/housing.csv')
```

3.1 Data exploration

```
[8]: housing.head()
```

```
[8]:  longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0    -122.23    37.88             41.0         880.0         129.0
1    -122.22    37.86             21.0        7099.0        1106.0
2    -122.24    37.85             52.0        1467.0         190.0
3    -122.25    37.85             52.0        1274.0         235.0
4    -122.25    37.85             52.0        1627.0         280.0

      population  households  median_income  median_house_value  ocean_proximity
0         322.0        126.0         8.3252         452600.0        NEAR BAY
1        2401.0       1138.0         8.3014        358500.0        NEAR BAY
2         496.0        177.0         7.2574        352100.0        NEAR BAY
3         558.0        219.0         5.6431        341300.0        NEAR BAY
4         565.0        259.0         3.8462        342200.0        NEAR BAY
```

Get information about all the columns

```
[55]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

Count the unique values in column (e.g. *ocean_proximity*)

```
[56]: housing['ocean_proximity'].value_counts()
```

```
[56]: <1H OCEAN      9136
      INLAND      6551
      NEAR OCEAN   2658
      NEAR BAY     2290
      ISLAND         5
      Name: ocean_proximity, dtype: int64
```

Summary statistics of the columns

```
[57]: housing.describe()
```

```
[57]:
```

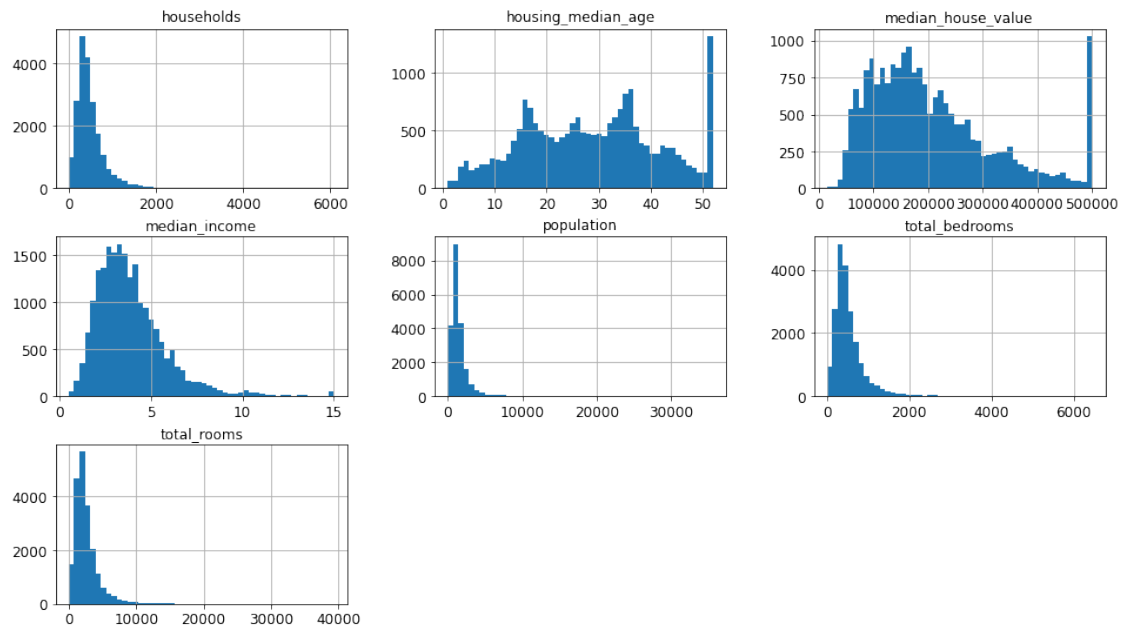
	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000


Simple visualization of the distribution of a subset of features: 'households', 'housing_median_age', 'median_house_value', 'median_income', 'population', 'total_bedrooms', 'total_rooms'

```
[58]: housing.hist(column=['households', 'housing_median_age',  
    ↪ 'median_house_value', 'median_income', 'population', 'total_bedrooms', 'total_rooms'],  
    bins=50,  
    figsize=(16,9))  
plt.show()
```

```
[59]: housing['median_income'].hist()
```

```
[59]: <AxesSubplot:>
```



```
[60]: housing['income_cat'] = pd.cut(housing['median_income'],
                                     bins=[0,1.5,3,4.5,6, np.inf],
                                     labels = ['Very Low', 'Low', 'Medium', 'High', 'Very High'])
housing
```

```
[60]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	
...	
20635	-121.09	39.48	25.0	1665.0	374.0	
20636	-121.21	39.49	18.0	697.0	150.0	
20637	-121.22	39.43	17.0	2254.0	485.0	
20638	-121.32	39.43	18.0	1860.0	409.0	
20639	-121.24	39.37	16.0	2785.0	616.0	

	population	households	median_income	median_house_value	\
0	322.0	126.0	8.3252	452600.0	
1	2401.0	1138.0	8.3014	358500.0	
2	496.0	177.0	7.2574	352100.0	
3	558.0	219.0	5.6431	341300.0	
4	565.0	259.0	3.8462	342200.0	
...	
20635	845.0	330.0	1.5603	78100.0	
20636	356.0	114.0	2.5568	77100.0	
20637	1007.0	433.0	1.7000	92300.0	
20638	741.0	349.0	1.8672	84700.0	
20639	1387.0	530.0	2.3886	89400.0	

	ocean_proximity	income_cat
0	NEAR BAY	Very High
1	NEAR BAY	Very High
2	NEAR BAY	Very High
3	NEAR BAY	High
4	NEAR BAY	Medium
...
20635	INLAND	Low
20636	INLAND	Low
20637	INLAND	Low
20638	INLAND	Low
20639	INLAND	Low

[20640 rows x 11 columns]

4 Partitioning the dataset into separate training and test sets

4.1 1) Random partition

```
[61]: from sklearn.model_selection import train_test_split

train_set, test_set = train_test_split(housing, test_size = 0.2, random_state = 42)
```

```
[62]: train_set.head()
```

```
[62]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
14196	-117.03	32.71	33.0	3126.0	627.0	
8267	-118.16	33.77	49.0	3382.0	787.0	
17445	-120.48	34.66	4.0	1897.0	331.0	
14265	-117.11	32.69	36.0	1421.0	367.0	
2271	-119.80	36.78	43.0	2382.0	431.0	

	population	households	median_income	median_house_value	\
14196	2300.0	623.0	3.2596	103000.0	
8267	1314.0	756.0	3.8125	382100.0	
17445	915.0	336.0	4.1563	172600.0	
14265	1418.0	355.0	1.9425	93400.0	
2271	874.0	380.0	3.5542	96500.0	

	ocean_proximity	income_cat
14196	NEAR OCEAN	Medium
8267	NEAR OCEAN	Medium
17445	NEAR OCEAN	Medium
14265	NEAR OCEAN	Low
2271	INLAND	Medium

We assign 20% of the sample to the test and the remaining 80% to the training, BUT no guarantee both training and test sets have the same label/outcome distribution, especially when the dataset is small. Let's see...

```
[63]: def income_cat_proportions(data):
      return data['income_cat'].value_counts() / len(data)
```

```
[64]: compare_props = pd.DataFrame({
      'Overall': income_cat_proportions(housing),
      'Random' : income_cat_proportions(test_set)
    }).sort_index()
```

```
[65]: compare_props['Rand %error'] = 100 * compare_props['Random'] /
      compare_props['Overall'] - 100
```

```
[67]: compare_props
```

```
[67]:
```

	Overall	Random	Rand %error
Very Low	0.039826	0.040213	0.973236
Low	0.318847	0.324370	1.732260
Medium	0.350581	0.358527	2.266446
High	0.176308	0.167393	-5.056334
Very High	0.114438	0.109496	-4.318374

4.2 2) Stratified Sampling

```
[68]: # StratifiedShuffleSplit creates splits by preserving the same percentage
# for each target class as in the complete set.
from sklearn.model_selection import StratifiedShuffleSplit

splitObject = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)

train_index, test_index = next(splitObject.split(housing,
→housing['income_cat']))

stratified_train_set = housing.loc[train_index]
stratified_test_set = housing.loc[test_index]

stratified_train_set.shape, stratified_test_set.shape, housing.shape
```

```
[68]: ((16512, 11), (4128, 11), (20640, 11))
```

```
[69]: compare_props['Stratified'] = stratified_test_set['income_cat'].value_counts() /
→ len(stratified_test_set)
```

```
[70]: compare_props['Stratified %error'] = 100 * compare_props['Stratified'] /
→ compare_props['Overall'] - 100
```

```
[71]: compare_props
```

```
[71]:
```

	Overall	Random	Rand %error	Stratified	Stratified %error
Very Low	0.039826	0.040213	0.973236	0.039729	-0.243309
Low	0.318847	0.324370	1.732260	0.318798	-0.015195
Medium	0.350581	0.358527	2.266446	0.350533	-0.013820
High	0.176308	0.167393	-5.056334	0.176357	0.027480
Very High	0.114438	0.109496	-4.318374	0.114583	0.127011

5 Prepare the data for Machine Learning algorithms

```
[72]: stratified_train_set
```

```
[72]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
16126	-122.47	37.79	52.0	437.0	105.0	
17709	-121.82	37.33	23.0	3279.0	647.0	
2501	-120.38	36.76	25.0	991.0	272.0	
2123	-119.71	36.76	28.0	2675.0	527.0	
2144	-119.76	36.77	36.0	2507.0	466.0	
...	
3382	-118.27	34.25	35.0	779.0	143.0	
841	-122.08	37.59	16.0	1816.0	365.0	
11749	-121.15	38.80	20.0	2104.0	370.0	
3940	-118.59	34.21	34.0	1943.0	320.0	
18827	-122.26	41.66	17.0	1885.0	350.0	

	population	households	median_income	median_house_value	\
16126	194.0	87.0	2.8125	500001.0	
17709	2582.0	630.0	4.3782	175800.0	
2501	941.0	262.0	1.8125	58000.0	
2123	1392.0	521.0	2.3108	72000.0	
2144	1227.0	474.0	2.7850	72300.0	
...	
3382	371.0	150.0	4.6635	230100.0	
841	1367.0	355.0	4.2350	156300.0	
11749	745.0	314.0	4.1685	217500.0	
3940	895.0	305.0	5.0462	227700.0	
18827	953.0	328.0	2.1607	61400.0	

	ocean_proximity	income_cat
16126	NEAR BAY	Low
17709	<1H OCEAN	Medium
2501	INLAND	Low
2123	INLAND	Low
2144	INLAND	Low
...
3382	<1H OCEAN	High
841	NEAR BAY	Medium
11749	INLAND	Medium
3940	<1H OCEAN	High
18827	INLAND	Low

[16512 rows x 11 columns]

```
[73]: housing = stratified_train_set.drop('median_house_value',axis=1)
housing_label = stratified_train_set['median_house_value'].copy()
housing.columns
```

```
[73]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
          'total_bedrooms', 'population', 'households', 'median_income',
```

```
'ocean_proximity', 'income_cat'],
dtype='object')
```

```
[74]: housing_label
```

```
[74]: 16126    500001.0
      17709    175800.0
      2501     58000.0
      2123     72000.0
      2144     72300.0
      ...
      3382    230100.0
      841     156300.0
      11749   217500.0
      3940    227700.0
      18827    61400.0
      Name: median_house_value, Length: 16512, dtype: float64
```

5.1 Identifying missing values

```
[76]: sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
      sample_incomplete_rows
```

```
[76]:   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
8383    -118.36    33.96                26.0        3543.0             NaN
10915   -117.87    33.73                45.0        2264.0             NaN
11311   -117.96    33.78                33.0        1520.0             NaN
696     -122.10    37.69                41.0         746.0             NaN
15137   -116.91    32.83                16.0        5203.0             NaN

      population  households  median_income  ocean_proximity  income_cat
8383         2742.0        951.0         2.5504      <1H OCEAN         Low
10915         1970.0        499.0         3.4193      <1H OCEAN       Medium
11311          658.0        242.0         4.8750      <1H OCEAN       High
696          387.0        161.0         3.9063      NEAR BAY       Medium
15137         2515.0        862.0         4.1050      <1H OCEAN       Medium
```

5.2 Eliminating rows with missing values

```
[78]: sample_incomplete_rows.dropna(subset=['total_bedrooms'], axis=0)
```

```
[78]: Empty DataFrame
      Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms,
      population, households, median_income, ocean_proximity, income_cat]
      Index: []
```

5.3 Eliminating variables with missing values

```
[79]: sample_incomplete_rows.dropna(axis=1)
```

```
[79]:
```

	longitude	latitude	housing_median_age	total_rooms	population	\
8383	-118.36	33.96	26.0	3543.0	2742.0	
10915	-117.87	33.73	45.0	2264.0	1970.0	
11311	-117.96	33.78	33.0	1520.0	658.0	
696	-122.10	37.69	41.0	746.0	387.0	
15137	-116.91	32.83	16.0	5203.0	2515.0	

	households	median_income	ocean_proximity	income_cat
8383	951.0	2.5504	<1H OCEAN	Low
10915	499.0	3.4193	<1H OCEAN	Medium
11311	242.0	4.8750	<1H OCEAN	High
696	161.0	3.9063	NEAR BAY	Medium
15137	862.0	4.1050	<1H OCEAN	Medium

5.4 Imputing missing values

5.4.1 1) Pandas

```
[80]: median = housing['total_bedrooms'].median()  
sample_incomplete_rows['total_bedrooms'].fillna(median, inplace=True)  
sample_incomplete_rows
```

```
[80]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
8383	-118.36	33.96	26.0	3543.0	437.0	
10915	-117.87	33.73	45.0	2264.0	437.0	
11311	-117.96	33.78	33.0	1520.0	437.0	
696	-122.10	37.69	41.0	746.0	437.0	
15137	-116.91	32.83	16.0	5203.0	437.0	

	population	households	median_income	ocean_proximity	income_cat
8383	2742.0	951.0	2.5504	<1H OCEAN	Low
10915	1970.0	499.0	3.4193	<1H OCEAN	Medium
11311	658.0	242.0	4.8750	<1H OCEAN	High
696	387.0	161.0	3.9063	NEAR BAY	Medium
15137	2515.0	862.0	4.1050	<1H OCEAN	Medium

5.4.2 2) Scikit-Learn

The **SimpleImputer** class.

The **SimpleImputer** class provides basic strategies for imputing missing values. Missing values can be imputed with a provided constant value, or using the statistics (mean, median or most frequent) of each column in which the missing values are located.

```
[81]: from sklearn.impute import SimpleImputer
      imputer = SimpleImputer(missing_values = np.nan, strategy = 'median')
```

Remove the text attribute because median can only be calculated on numerical attributes:

```
[36]: housing_num = housing.select_dtypes(include=[np.number])
```

```
[37]: imputer.fit(housing_num)
```

```
[37]: SimpleImputer(strategy='median')
```

```
[38]: imputer.statistics_
```

```
[38]: array([[-118.5    ,  34.26   ,  29.    , 2137.    ,  437.    , 1170.    ,
           411.    ,   3.5375])
```

Transform the training set:

```
[39]: X = imputer.transform(housing_num)
      X
```

```
[39]: array([[ -1.2247e+02,  3.7790e+01,  5.2000e+01, ...,  1.9400e+02,
           8.7000e+01,  2.8125e+00],
          [ -1.2182e+02,  3.7330e+01,  2.3000e+01, ...,  2.5820e+03,
           6.3000e+02,  4.3782e+00],
          [ -1.2038e+02,  3.6760e+01,  2.5000e+01, ...,  9.4100e+02,
           2.6200e+02,  1.8125e+00],
          ...,
          [ -1.2115e+02,  3.8800e+01,  2.0000e+01, ...,  7.4500e+02,
           3.1400e+02,  4.1685e+00],
          [ -1.1859e+02,  3.4210e+01,  3.4000e+01, ...,  8.9500e+02,
           3.0500e+02,  5.0462e+00],
          [ -1.2226e+02,  4.1660e+01,  1.7000e+01, ...,  9.5300e+02,
           3.2800e+02,  2.1607e+00]])
```

Scikit-Learn API is organized around a bunch of design principles:

Consistency: all object share a consistent and simple interface

Estimator: object that can estimate some parameters. Estimation performed by the method fit which takes only a dataset as parameter, any other parameter is an hyperparameter

Transformers: some estimators transform a dataset. The transformation is performed by the method transform with the dataset to transform as a parameter. It returns the transformed dataset. There is a convenient fit_transform method, which is optimized and runs much faster

Predictors: some estimator are able to make predictions. A predictor has a method predict that takes a dataset of new samples and returns the corresponding predictions

Inspection: all hyperparameter are accessible via instance variable as well as the

- Nonproliferation of classes: datasets are Numpy arrays or Scipy sparse matrices. No
- Composition: existing building block are reusable
- Sensible defaults: reasonable default values.

```
[82]: imputer.strategy
```

```
[82]: 'median'
```

```
[41]: housing_trasformed = pd.DataFrame(X, columns= housing_num.columns, index =  
→list(housing.index.values))
```

5.5 Encoding nominal features

```
[87]: housing_cat = housing[['ocean_proximity']]  
housing_cat
```

```
[87]:      ocean_proximity  
16126      NEAR BAY  
17709      <1H OCEAN  
2501       INLAND  
2123       INLAND  
2144       INLAND  
...      ...  
3382      <1H OCEAN  
841       NEAR BAY  
11749     INLAND  
3940      <1H OCEAN  
18827     INLAND  
  
[16512 rows x 1 columns]
```

To convert categorical features to such integer codes, we can use the **OrdinalEncoder**. This estimator transforms each categorical feature to one new feature of integers (0 to n_categories - 1)

```
[88]: from sklearn.preprocessing import OrdinalEncoder
```

```
[89]: ordinal_encoder = OrdinalEncoder()  
housing_cat_encoded = ordinal_encoder.fit_transform(housing_cat)  
housing_cat_encoded[:10]
```

```
[89]: array([[3.],  
          [0.],  
          [1.],  
          [1.],  
          [1.],  
          [0.],  
          [4.]])
```

```
[1.],  
[0.],  
[1.]])
```

Such integer representation can, however, not be used directly with all scikit-learn estimators, as these expect continuous input, and would interpret the categories as being ordered, which is often not desired. A common workaround to this issue is to use a technique called **one-hot encoding**

```
[90]: from sklearn.preprocessing import OneHotEncoder  
  
cat_encoder = OneHotEncoder()  
housing_cat_1hot = cat_encoder.fit_transform(housing_cat)
```

By default, the `OneHotEncoder` class returns a sparse array, but we can convert it to a dense array if needed by calling the `toarray()` method:

```
[91]: housing_cat_1hot.toarray()[0:6]
```

```
[91]: array([[0., 0., 0., 1., 0.],  
          [1., 0., 0., 0., 0.],  
          [0., 1., 0., 0., 0.],  
          [0., 1., 0., 0., 0.],  
          [0., 1., 0., 0., 0.],  
          [1., 0., 0., 0., 0.]])
```

Alternatively, you can set `sparse=False` when creating the `OneHotEncoder`:

```
[56]: cat_encoder = OneHotEncoder(sparse = False)  
housing_cat_1hot = cat_encoder.fit_transform(housing_cat)  
type(housing_cat_1hot)
```

```
[56]: numpy.ndarray
```

```
[57]: cat_encoder.categories_
```

```
[57]: [array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],  
        dtype=object)]
```

5.6 Attributes creation

Let's create a new transformer to add extra attributes. All you need is to convert an existing Python function into a transformer to assist in data cleaning or processing. You can implement a transformer from an arbitrary function with the class **FunctionTransformer**

```
[92]: housing.columns
```

```
[92]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',  
        'total_bedrooms', 'population', 'households', 'median_income',
```

```
    'ocean_proximity', 'income_cat'],
    dtype='object')
```

```
[100]: rooms_ix, bed_rooms_ix, population_ix, household_ix = [
        list(housing.columns).index(col) for col in
        ↪ ['total_rooms', 'total_bedrooms', 'population', 'households']
    ]

    def add_extra_features(X):
        roomsXhouse = X[:, rooms_ix] / X[:, household_ix]
        popXhouse = X[:, population_ix] / X[:, household_ix]
        return np.c_[X, roomsXhouse, popXhouse]

    from sklearn.preprocessing import FunctionTransformer
    attr_adder = FunctionTransformer(add_extra_features, validate = False)

    housinhg_extra = attr_adder.fit_transform(housing.values)
```

```
[101]: housinhg_extra_df = pd.DataFrame(housinhg_extra,
                                         columns = list(housing.columns)+['a', 'b'])
    housinhg_extra_df.head()
```

```
[101]: longitude latitude housing_median_age total_rooms total_bedrooms population \
0    -122.47    37.79          52.0         437.0         105.0        194.0
1    -121.82    37.33          23.0        3279.0         647.0       2582.0
2    -120.38    36.76          25.0         991.0         272.0        941.0
3    -119.71    36.76          28.0        2675.0         527.0       1392.0
4    -119.76    36.77          36.0        2507.0         466.0       1227.0

    households median_income ocean_proximity income_cat      a      b
0         87.0         2.8125      NEAR BAY      Low  5.022989  2.229885
1        630.0         4.3782    <1H OCEAN    Medium  5.204762  4.098413
2        262.0         1.8125      INLAND      Low  3.782443  3.591603
3        521.0         2.3108      INLAND      Low  5.134357  2.671785
4        474.0         2.785      INLAND      Low  5.28903  2.588608
```

5.7 Attribute or feature scaling

ML algorithms don't perform well when the numerical attributes have very different scales. Two classes to report all the attributes to the same scale:

- **Mix-max scaling:** SkLearn provides the transformer **MinMaxScaler**
- **Standardization:** SkLearn provides the transformer **StandardScaler**

5.8 Transformation Pipeline

Since there are many transformation steps that need to be executed in the right order, need a way to automatically create this sequence of transformation. SkLearn provides the **Pipeline** class. This

class takes an arbitrary number of SkLearn transformers, as a list of name/estimator pairs. When you call the method *fit()*, it runs the method *fit_transform()* of each element in list, sequentially

Now let's build a pipeline for preprocessing the numerical attributes:

```
[103]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy = 'median')),
    ('attrs_adder', FunctionTransformer(add_extra_features, validate=False)),
    ('std_scaler', StandardScaler())
])

housing_num_tr = num_pipeline.fit_transform(housinhg_extra_df)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-103-b29347a14477> in <module>
      8 ]
      9
----> 10 housing_num_tr = num_pipeline.fit_transform(housinhg_extra_df)

/opt/anaconda3/lib/python3.8/site-packages/sklearn/pipeline.py in
-> fit_transform(self, X, y, **fit_params)
    376     """
    377     fit_params_steps = self._check_fit_params(**fit_params)
--> 378     Xt = self._fit(X, y, **fit_params_steps)
    379
    380     last_step = self._final_estimator

/opt/anaconda3/lib/python3.8/site-packages/sklearn/pipeline.py in _fit(self, X,
-> y, **fit_params_steps)
    301         cloned_transformer = clone(transformer)
    302         # Fit or load from cache the current transformer
--> 303         X, fitted_transformer = fit_transform_one_cached(
    304             cloned_transformer, X, y, None,
    305             message_clsname='Pipeline',

/opt/anaconda3/lib/python3.8/site-packages/joblib/memory.py in __call__(self,
-> *args, **kwargs)
    350
    351     def __call__(self, *args, **kwargs):
--> 352         return self.func(*args, **kwargs)
    353
    354     def call_and_shelve(self, *args, **kwargs):
```

```

/opt/anaconda3/lib/python3.8/site-packages/sklearn/pipeline.py in
↳ _fit_transform_one(transformer, X, y, weight, message_clsname, message,
↳ **fit_params)
    752     with _print_elapsed_time(message_clsname, message):
    753         if hasattr(transformer, 'fit_transform'):
--> 754             res = transformer.fit_transform(X, y, **fit_params)
    755         else:
    756             res = transformer.fit(X, y, **fit_params).transform(X)

/opt/anaconda3/lib/python3.8/site-packages/sklearn/base.py in
↳ fit_transform(self, X, y, **fit_params)
    697         if y is None:
    698             # fit method of arity 1 (unsupervised transformation)
--> 699             return self.fit(X, **fit_params).transform(X)
    700         else:
    701             # fit method of arity 2 (supervised transformation)

/opt/anaconda3/lib/python3.8/site-packages/sklearn/impute/_base.py in fit(self,
↳ X, y)
    286         self : SimpleImputer
    287         """
--> 288         X = self._validate_input(X, in_fit=True)
    289
    290         # default fill_value is 0 for numerical input and "missing_valu

/opt/anaconda3/lib/python3.8/site-packages/sklearn/impute/_base.py in
↳ _validate_input(self, X, in_fit)
    258             new_ve = ValueError("Cannot use {} strategy with
↳
↳ non-numeric "
    259                                     "data:\n{}".format(self.strategy,
↳
↳ ve))
--> 260             raise new_ve from None
    261         else:
    262             raise ve

ValueError: Cannot use median strategy with non-numeric data:
could not convert string to float: 'NEAR BAY'

```

```
[62]: housing_num_tr.shape
```

```
[62]: (16512, 10)
```

If you have a Pandas DataFrame it is now preferable to use the **ColumnTransformer** class that was introduced in SkLearn 0.20.

```
[104]: from sklearn.compose import ColumnTransformer
```

```
[106]: num_attribs = list(housing_num)
cat_attribs = ['ocean_proximity']

full_pipeline = ColumnTransformer([
    ('num', num_pipeline, num_attribs),
    ('cat', OneHotEncoder(), cat_attribs)
])

housing_final = full_pipeline.fit_transform(housing)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-106-1a782cbf41f5> in <module>
----> 1 num_attribs = list(housing_num)
      2 cat_attribs = ['ocean_proximity']
      3
      4 full_pipeline = ColumnTransformer([
      5     ('num', num_pipeline, num_attribs),

NameError: name 'housing_num' is not defined
```

```
[65]: housing_final.shape, housing.values.shape
```

```
[65]: ((16512, 15), (16512, 10))
```

```
[68]: housing_final[0:6]
```

```
[68]: array([[ -1.44853942,  1.00749903,  1.85042332, -1.00225667, -1.02608971,
        -1.09987832, -1.0732447 , -0.55467031, -0.17413103, -0.07471997,
         0.          ,  0.          ,  0.          ,  1.          ,  0.          ],
       [-1.12372271,  0.79233009, -0.44832649,  0.28300714,  0.2522925 ,
         1.0191734 ,  0.32798234,  0.26444129, -0.09511204,  0.08667068,
         1.          ,  0.          ,  0.          ,  0.          ,  0.          ],
       [-0.40412878,  0.52570771, -0.28979202, -0.75171615, -0.63219704,
        -0.43700912, -0.62165219, -1.0778303 , -0.71341051,  0.04289592,
         0.          ,  1.          ,  0.          ,  0.          ,  0.          ],
       [-0.06931772,  0.52570771, -0.05199032,  0.00985466, -0.03074415,
        -0.03680296,  0.04670472, -0.81713968, -0.12571786, -0.03655167,
         0.          ,  1.          ,  0.          ,  0.          ,  0.          ],
       [-0.09430362,  0.5303853 ,  0.58214756, -0.06612152, -0.17462112,
        -0.18321985, -0.07458012, -0.56905721, -0.05847994, -0.04373597,
         0.          ,  1.          ,  0.          ,  0.          ,  0.          ],
       [ 0.72023674, -0.84950247,  1.21628544, -0.46770993, -0.61332793,
        -0.54881839, -0.56488056,  0.87481206,  0.12047978, -0.01945555,
         1.          ,  0.          ,  0.          ,  0.          ,  0.          ]])
```

6 Extra material

6.1 Model persistence using joblib

```
[105]: my_model = full_pipeline
```

```
-----  
NameError                                Traceback (most recent call last)  
<ipython-input-105-c829ee97a232> in <module>  
----> 1 my_model = full_pipeline  
  
NameError: name 'full_pipeline' is not defined
```

```
[71]: #from sklearn.externals import joblib  
import joblib  
joblib.dump(my_model, "full_pipeline.pkl") # DIFF  
my_model_loaded = joblib.load("full_pipeline.pkl") # DIFF
```

```
[72]: my_model_loaded
```

```
[72]: ColumnTransformer(transformers=[('num',  
                                     Pipeline(steps=[('imputer',  
SimpleImputer(strategy='median')),  
                                     ('attrs_adder',  
FunctionTransformer(func=<function add_extra_features at 0x7fe58e0aaa60>)),  
                                     ('std_scaler',  
                                     StandardScaler())]),  
                                ['longitude', 'latitude', 'housing_median_age',  
                                'total_rooms', 'total_bedrooms', 'population',  
                                'households', 'median_income']),  
                                ('cat', OneHotEncoder(), ['ocean_proximity'])])
```

6.2 Some further examples of using pickle

```
[107]: import pandas as pd  
import pickle
```

```
[109]: print ('convert: csv -> pkl')  
datimaggio18 = pd.read_csv('01_2018.csv', delimiter=';', header=0,  
    ↳ encoding='mac_roman')  
datimaggio18
```

convert: csv -> pkl

```
[109]:      Bicicletta Tipo_bici  Cliente Data_riferimento_prelievo \  
0      7486      Bike      141116      01/01/18
```

1	8279	Bike	265468	01/01/18
2	1284	Bike	232605	01/01/18
3	7411	Bike	21489	01/01/18
4	1730	Bike	220370	01/01/18
...
250156	7780	Bike	308325	31/01/18
250157	3562	Bike	163545	31/01/18
250158	11108	eBike	81098	31/01/18
250159	7828	Bike	17302	31/01/18
250160	6992	Bike	234044	31/01/18

	Data_prelievo	Ora_prelievo	Giorno_prelievo	Mese_prelievo	\
0	01/01/18 07:18	7	1	1	
1	01/01/18 07:35	7	1	1	
2	01/01/18 07:49	7	1	1	
3	01/01/18 07:56	7	1	1	
4	01/01/18 07:58	7	1	1	
...	
250156	01/02/18 00:37	0	1	2	
250157	01/02/18 00:42	0	1	2	
250158	01/02/18 00:52	0	1	2	
250159	01/02/18 00:58	0	1	2	
250160	01/02/18 00:59	0	1	2	

	Anno_prelievo	GMA_prelievo	...	Precipitazioni_GA	Press_atm_GA	\
0	2018	112018	...	Yes	1012	
1	2018	112018	...	Yes	1012	
2	2018	112018	...	Yes	1012	
3	2018	112018	...	Yes	1012	
4	2018	112018	...	Yes	1012	
...	
250156	2018	122018	...	Yes	1005	
250157	2018	122018	...	Yes	1005	
250158	2018	122018	...	Yes	1005	
250159	2018	122018	...	Yes	1005	
250160	2018	122018	...	Yes	1005	

	Pm10_GP	Pm25_GP	No2_mean_GP	Co_mean_GP	Pm10_GA	Pm25_GA	No2_mean_GA	\
0	39	38	40.1	0.9	39	38	40.1	
1	39	38	40.1	0.9	39	38	40.1	
2	39	38	40.1	0.9	39	38	40.1	
3	39	38	40.1	0.9	39	38	40.1	
4	39	38	40.1	0.9	39	38	40.1	
...	
250156	35	29	44.0	0.7	35	29	44.0	
250157	35	29	44.0	0.7	35	29	44.0	
250158	35	29	44.0	0.7	35	29	44.0	

250159	35	29	44.0	0.7	35	29	44.0
250160	35	29	44.0	0.7	35	29	44.0

	Co_mean_GA
0	0.9
1	0.9
2	0.9
3	0.9
4	0.9
...	...
250156	0.7
250157	0.7
250158	0.7
250159	0.7
250160	0.7

[250161 rows x 52 columns]

```
[110]: pickle.dump(datimaggio18, open("datimaggio18.pkl", "wb"))
```

```
[111]: bikemi = pd.read_pickle('datimaggio18.pkl'.format(5,2018))
```

```
[78]: print('Number of rents')
      len(bikemi)
```

Number of rents

```
[78]: 250161
```

```
[79]: bikemi['Cliente'].nunique()
```

```
[79]: 24944
```

7 Let's see a bit of statistical visualization tools

```
[113]: from IPython.core.pylabtools import figsize
      import numpy as np
      from matplotlib import pyplot as plt
      figsize(12.5, 4)


      import scipy.stats as stats
      a = np.arange(16)
      poi = stats.poisson
      lambda_ = [1.5, 4.25]
      colours = ["red", "blue"]
```

```
plt.bar(a, poi.pmf(a, lambda_[0]), color=colours[0],
        label="$\\lambda = %.1f$" % lambda_[0], alpha=0.60,
        edgecolor=colours[0], lw="3")

plt.bar(a, poi.pmf(a, lambda_[1]), color=colours[1],
        label="$\\lambda = %.1f$" % lambda_[1], alpha=0.60,
        edgecolor=colours[1], lw="3")

plt.xticks(a + 0.4, a)
plt.legend()
plt.ylabel("probability of $k$")
plt.xlabel("$k$")
plt.title("Probability mass function of a Poisson random variable with
↳different \\
$\\lambda$ values")
```

[113]: `Text(0.5, 1.0, 'Probability mass function of a Poisson random variable with different $\\lambda$ values')`

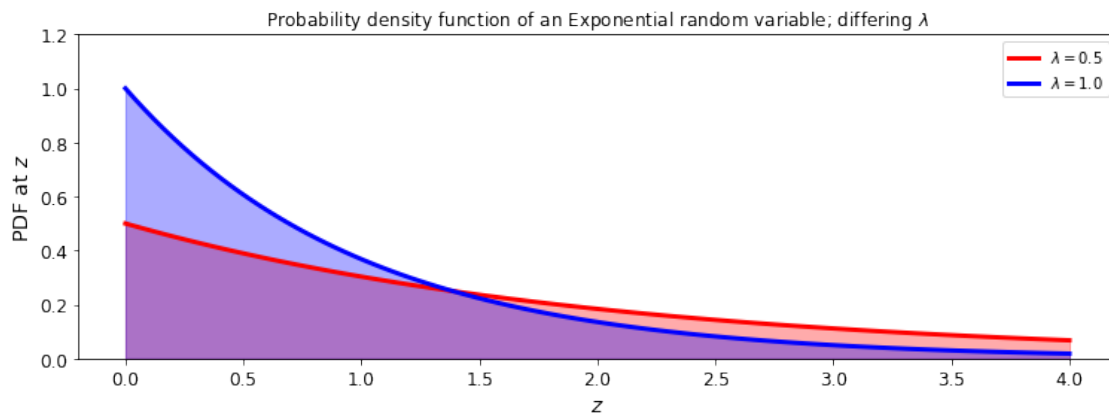


```
[114]: a = np.linspace(0, 4, 100)
expo = stats.expon
lambda_ = [0.5, 1]

for l, c in zip(lambda_, colours):
    plt.plot(a, expo.pdf(a, scale=1. / l), lw=3,
             color=c, label="$\\lambda = %.1f$" % l)
    plt.fill_between(a, expo.pdf(a, scale=1. / l), color=c, alpha=.33)

plt.legend()
plt.ylabel("PDF at $z$")
plt.xlabel("$z$")
plt.ylim(0, 1.2)
```

```
plt.title("Probability density function of an Exponential random variable;\ndiffering  $\lambda$ ");
```



```
[82]: figsize(12.5, 3.5)
count_data = np.loadtxt("/Users/giancarlomanzi/Documents/Box Sync BackUp PC\
↳Lavoro 24062015/documenti/Ricerca/ALTERNATIVE ERASMUS PROJECT/NUOVO PROGETTO\
↳DI VISITING/Lectures/Topic 2 - Introduction to Python and the\
↳Anaconda-Jupyter environment - 3 hours/txtdata.csv")
n_count_data = len(count_data)
plt.bar(np.arange(n_count_data), count_data, color="#348ABD")
plt.xlabel("Time (days)")
plt.ylabel("count of text-msgs received")
plt.title("Did the user's texting habits change over time?")
plt.xlim(0, n_count_data);
```

```
-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-82-a439bff65e7e> in <module>
      1 figsize(12.5, 3.5)
----> 2 count_data = np.loadtxt("/Users/giancarlomanzi/Documents/Box Sync BackUp\
↳PC Lavoro 24062015/documenti/Ricerca/ALTERNATIVE ERASMUS PROJECT/NUOVO\
↳PROGETTO DI VISITING/Lectures/Topic 2 - Introduction to Python and the\
↳Anaconda-Jupyter environment - 3 hours/txtdata.csv")
      3 n_count_data = len(count_data)
      4 plt.bar(np.arange(n_count_data), count_data, color="#348ABD")
      5 plt.xlabel("Time (days)")

/opt/anaconda3/lib/python3.8/site-packages/numpy/lib/npymio.py in loadtxt(fname,
↳dtype, comments, delimiter, converters, skiprows, usecols, unpack, ndmin,
↳encoding, max_rows, like)
    1040         fname = os.fspath(fname)
    1041         if _is_string_like(fname):
-> 1042             fh = np.lib._datasource.open(fname, 'rt', encoding=encoding
```

```

1043         fencoding = getattr(fh, 'encoding', 'latin1')
1044         line_iter = iter(fh)

/opt/anaconda3/lib/python3.8/site-packages/numpy/lib/_datasource.py in
↳ open(path, mode, destpath, encoding, newline)
    191
    192     ds = DataSource(destpath)
--> 193     return ds.open(path, mode, encoding=encoding, newline=newline)
    194
    195

/opt/anaconda3/lib/python3.8/site-packages/numpy/lib/_datasource.py in
↳ open(self, path, mode, encoding, newline)
    530                                     encoding=encoding, newline=newlin )
    531     else:
--> 532         raise FileNotFoundError(f"{path} not found.")
    533
    534

FileNotFoundError: /Users/giancarlomanzi/Documents/Box Sync BackUp PC Lavoro
↳ 24062015/documenti/Ricerca/ALTERNATIVE ERASMUS PROJECT/NUOVO PROGETTO DI
↳ VISITING/Lectures/Topic 2 - Introduction to Python and the Anaconda-Jupyter
↳ environment - 3 hours/txtdata.csv not found.

```

7.1 Pipelines in text mining/natural language processing

- We spend a lot of time in pre-processing and cleaning data
- Therefore we need to create multipurpose software objects to be used in different situations.
- For this we can use the Pipeline tool in scikit-learn.
- It is composed by *transformers* (tools for transforming data, for example to normalize a variable) and *estimators* (for example a fitting or predicting tool).
- All transformers and estimators in scikit-learn are implemented as Python *classes*, each with their own attributes and methods.
- We use *inherited* classes from scikit-learn to implement our own class.

```

[115]: # Example of inheritance
from sklearn.preprocessing import OneHotEncoder
#Some data:
X = [[0, 0, 3], [1, 1, 0], [0, 2, 1], [1, 0, 2]]
#Initializing an object of class OneHotEncoder
# Here we are "inheriting" classes from OneHotEncoder into the variable
↳ 'one_hot_enc'
one_hot_enc = OneHotEncoder( sparse = True )

#Calling methods on our OneHotEncoder object
one_hot_enc.fit( X ) #returns nothing
transformed_data = one_hot_enc.transform(X).toarray() #returns something

```

```
#fit_transformed_data = one_hot_enc.transform( X ) #returns something
```

```
[116]: print(pd.DataFrame(X))  
#Comments: The first column takes on 2 values, the second 3 and the fourth 4
```

```
    0  1  2  
0  0  0  3  
1  1  1  0  
2  0  2  1  
3  1  0  2
```

```
[117]: print(transformed_data)  
#Comments: the first two columns express the binary coding of the first  
↪ "feature";  
# the next three columns express the binary coding of the second "feature";  
# The next four columns express the binary coding of the third "feature";
```

```
[[1.  0.  1.  0.  0.  0.  0.  0.  1.]  
 [0.  1.  0.  1.  0.  1.  0.  0.  0.]  
 [1.  0.  0.  0.  1.  0.  1.  0.  0.]  
 [0.  1.  1.  0.  0.  0.  0.  1.  0.]]
```

7.2 Pipelines in text mining/natural language processing (2)

- Our own transformer will be formed by inheriting from some other scikit-learn class.
- See a tutorial here <https://www.programiz.com/python-programming/class> about classes and objects in python and a tutorial here <https://www.programiz.com/python-programming/inheritance> about inheritance.
- The base classes inherited from scikit-learn are TransformerMixin (<https://scikit-learn.org/stable/modules/generated/sklearn.base.TransformerMixin.html>) and BaseEstimator (<https://scikit-learn.org/stable/modules/generated/sklearn.base.BaseEstimator.html>).

```
[118]: import numpy as np  
import pandas as pd
```

```
[121]: # Load Data  
pd.set_option('display.max_colwidth', None)  
BikeSent = pd.read_csv("BikeMiSentiment2019_UTF-8.csv", sep=';')  
BikeSent
```

```
[121]:      v1 \  
0      positive  
1      positive  
2      negative  
3      positive  
4      positive  
..      ...  
995    negative
```

996 positive
997 positive
998 positive
999 positive

v2

0 When I had problems with the return of the bike, the assistance was very kind, but could not solve the problem quickly nor prevent me from being charged for the rental because I had exceeded half an hour (not because of the route I took carried out, but due to the impossibility of hanging up the bike in 2 different stations). It would be desirable to be able to solve problems more efficiently or at least not to charge the user in the event of a reported malfunction. For the rest, when the service works, it's really practical and useful.\t

1
more electric bikes. often even if present they are not available when there are few, why?\t

2
pay more attention to stations that very often are without bicycles or full and do not allow their repositioning\t

3
essential to insert bikes with child seats\t

4
extension completed at train and metro stations not yet served\t

..

...

995
Main problem I think is the maintenance of traditional bikes, often you are forced to change bikes several times before finding a functioning one\t

996
I feel good but without a credit card you can't even buy a day card, it doesn't seem right because students like me often only have a prepaid card\t

997
I don't have any suggestions at the moment. the comment, thank you for the excellent service provided.\t

998
I would like it if the number of red ebikes increased considerably\t
999
need more maintenance, stations in the center with too many bikes\t

[1000 rows x 2 columns]

```
[122]: # Rename columns
BikeSent.columns = ["target", "text"]
BikeSent.head()
```

```
[122]:      target \
0  positive
1  positive
2  negative
3  positive
4  positive
```

text

```
0  When I had problems with the return of the bike, the assistance was very
kind, but could not solve the problem quickly nor prevent me from being charged
for the rental because I had exceeded half an hour (not because of the route I
took carried out, but due to the impossibility of hanging up the bike in 2
different stations). It would be desirable to be able to solve problems more
efficiently or at least not to charge the user in the event of a reported
malfunction. For the rest, when the service works, it's really practical and
useful.\t
1
more electric bikes. often even if present they are not available when there are
few, why?\t
2
pay more attention to stations that very often are without bicycles or full and
do not allow their repositioning\t
3
essential to insert bikes with child seats\t
4
extension completed at train and metro stations not yet served\t
```

```
[123]: # Encode categories
BikeSent['target'] = np.where(BikeSent['target']=='positive',1,0)
BikeSent.head()
```

```
[123]:      target \
0         1
1         1
2         0
3         1
4         1
```

text

```
0  When I had problems with the return of the bike, the assistance was very
kind, but could not solve the problem quickly nor prevent me from being charged
for the rental because I had exceeded half an hour (not because of the route I
took carried out, but due to the impossibility of hanging up the bike in 2
different stations). It would be desirable to be able to solve problems more
efficiently or at least not to charge the user in the event of a reported
malfunction. For the rest, when the service works, it's really practical and
useful.\t
```

```

1
more electric bikes. often even if present they are not available when there are
few, why?\t
2
pay more attention to stations that very often are without bicycles or full and
do not allow their repositioning\t
3
essential to insert bikes with child seats\t
4
extension completed at train and metro stations not yet served\t

```

```

[124]: # split the sample in train (used also for cross-validation) + test
from sklearn.model_selection import train_test_split
X = BikeSent[['text']]
y = BikeSent['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10,
↳random_state=42)

```

```

[125]: X_train

```

```

[125]:
text
716 Luckily I have tried the new bicycle models a few times and they are
definitely uncomfortable. I think there is a design error because the saddle is
too far forward and you have difficulty pedaling. I hope you realize this before
increasing the number of bikes you buy\t
351
Improve the bike pickup and storage system. For older people they are too heavy
to lift\t
936
I kindly ask you to make the stall n. 151 Balilla
- Tibaldi. Sometimes it is uninhabitable and there are few bicycles available or
they are generally few or poorly functioning (eg deflated wheels, poorly
functioning brakes, gearshift changes).\t
256
bikes should be maintained much, much better, often with badly maintained
bicycles and without brakes or even for the electrics that the battery does not
work\t
635
Increase maintenance\t
..
...
106
A really
useful service, I hope in the possibility of using 24 hours a day, especially
for us young people it can be very useful at night when the vehicles are almost
zero and you are forced to use taxis.\t
270
only problem to report too often the stalls do not record the correct
establishment of the bike and you risk icorrerere nela penalty\t

```


860 Some
discounts for the renewal of the subscription. The offers seem to me always and
only for the new subscribers. In addition, a few more conventions for Bikemi
subscribers who give discounts elsewhere.\t
435
the service is smart but is very limited by the location of the stations. They
are all a center. There isn't one in Stazione Lambrate or the eastern suburbs.\t
102
good\t

[900 rows x 1 columns]

7.3 Custom Transformers

7.3.1 Cleaning Text

- We create here our own transformer (which will be a class) inheriting the TransformerMixin and the BaseEstimator classes from scikit-learn

```
[126]: from sklearn.base import BaseEstimator
from sklearn.base import TransformerMixin
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import SnowballStemmer

# Custom Transformer (Inheriting from classes)
class CleanText( BaseEstimator, TransformerMixin ):

    # Class Constructor
    # The class constructor is formed by a function with double underscore __ :
    # these are called 'special functions' as they have special meaning.
    # In particular the '__init__' gets called whenever
    # a new object of that class is instantiated,
    # and are used to initialize all the necessary variables.
    # In this example we initialize the language variable 'lang' with 'English'
    # and pick the SnowballStemmer as the default stemmer.
    def __init__( self, lang = "english"):
        self.lang = lang
        self.stemmer = SnowballStemmer(self.lang)

    # The 'fit' method here is used to instantiate the class on the 'self'
    ↪variable
    # and return the object itself
    def fit( self, X, y = None ):
        return self

    # Custom function: this applies the stemmer just created in the '__init__'
    # part to the 'self' variable
```

```

def clean( self, x ):
    words = [self.stemmer.stem(word) for word in word_tokenize(x.lower())]
    →if word.isalpha() and word not in stopwords.words("english"):
        return " ".join(words)

    # Method that describes what we need this transformer to do i.e. cleaning
    →the text
    # in the 'text' column in the data frame.
    # This will be used later on in the usage of the custom transformer
    # within the pipeline.
def transform( self, X, y = None ):
    return X["text"].apply(self.clean)

```

7.3.2 Feature extraction

```

[127]: # Custom Transformer: same parts as the previous custom transformer
        # This one will be used for feature extraction

class CustomFeatures( BaseEstimator, TransformerMixin ):

    # Class Constructor
    def __init__( self ):
        return

    # Return self nothing else to do here
    def fit( self, X, y = None ):
        return self

    # Method that describes what we need this transformer to do i.e.
    # returning length, digits and punctuations in the 'text' column in data
    →frame
    def transform( self, X, y = None ):
        f = pd.DataFrame()
        f['len'] = X['text'].str.len()
        f['digits'] = X['text'].str.findall(r'\d').str.len()
        f['punct'] = X['text'].str.findall(r'^a-zA-Z\d\s:').str.len()
        return f[['len', 'digits', 'punct']]

```

7.4 Pipeline usage

7.4.1 Pipeline for data pre processing

```

[129]: from sklearn.pipeline import Pipeline
        from sklearn.pipeline import FeatureUnion
        # FeatureUnion combines two or more pipelines or transformers
        # and is very fast!
        from sklearn.feature_extraction.text import TfidfVectorizer

```

```

from sklearn.feature_selection import SelectKBest, chi2
from sklearn.preprocessing import StandardScaler
# Our first pipeline called 'pipe' will be formed by three 'steps' or parts:
# 1) "extract" which in turns is formed through FeatureUnion which
# put together two parts:
# "terms" (formed by a pipeline with the CleanText() transformer we created
→ above
# and the TfidfVectorizer text vectorizing transformer from scikit-learn) and
→ "custom"
# (formed by the CustomFeatures transformer we created above);
# 2) "select", formed by the scikit-learn transformer method "SelectKBest" for
→ feature
# selection with a chi squared score function;
# 3) "scale", same as 2) using the StandardScaler method from scikit-learn.
# The whole pipeline will be used as pre-processing task in classifying
→ pipelines.
pipe = Pipeline([("extract", FeatureUnion([("terms", Pipeline([('clean',
→ CleanText()),

                                                                    ('tfidf',
→ TfidfVectorizer()))]),

                                                                    ("custom", CustomFeatures()))]),
                ("select", SelectKBest(score_func = chi2)),
                ("scale", StandardScaler(with_mean = False))])

```

7.4.2 Classifier implemented through pipelines: Logistic Model

```

[131]: # Logistic Model
from sklearn.linear_model import LogisticRegression
pipe_logistic = Pipeline([('pre_process', pipe),
                           ('classify', LogisticRegression(max_iter=10000, tol=0.
→ 1, solver='lbfgs'))])

```

```

[132]: # Fit on training
pipe_logistic.fit(X_train, y_train)

```

```

[132]: Pipeline(steps=[('pre_process',
                        Pipeline(steps=[('extract',
                                         FeatureUnion(transformer_list=[('terms',
Pipeline(steps=[('clean',
                  CleanText()),
                  ('tfidf',
                    TfidfVectorizer()))]),
                                         ('custom',
CustomFeatures()))]),
                        ('select',
                          SelectKBest(score_func=<function chi2 at

```

```
0x7f83b3de78b0>)),
                                ('scale', StandardScaler(with_mean=False))])),
                                ('classify', LogisticRegression(max_iter=10000, tol=0.1)))
```

```
[133]: # Evaluate on test
# The F1 score can be interpreted as a weighted average of the precision and
→recall,
# where an F1 score reaches its best value at 1 and worst score at 0.
#The relative contribution of precision
# and recall to the F1 score are equal. The formula for the F1 score is:
#  $F1 = 2 * (precision * recall) / (precision + recall)$ 
from sklearn.metrics import f1_score
y_pred = pipe_logistic.predict(X_test)
f1_score(y_test, y_pred)
```

```
[133]: 0.7972972972972973
```

```
[98]: # we can classify new messages!
msg = pd.DataFrame(columns = ["text"],
                        data    = ["The bikes are heavy and unwieldy. The collection
→and return of the bicycle is super-comfortable because the bikes are heavy"])

pipe_logistic.predict(msg)
```

```
[98]: array([1])
```

```
[99]: # we can classify new messages!
#msg = pd.DataFrame(columns = ["text"],
#                        #data    = ["REMINDER FROM 02: To get 2.50 pounds free call
→credit and details of great offers pls reply 2 this text with your valid
→name, house no and postcode"])

msg = pd.DataFrame(columns = ["text"],
                    data    = ["Satisfied"])

pipe_logistic.predict(msg)
```

```
[99]: array([1])
```

7.5 Using bi-grams

```
[100]: # extract features
pipe_extract = FeatureUnion([("terms", Pipeline([('clean', CleanText()),
                                                    ('tfidf',
→TfidfVectorizer())])),
```

```
("custom", CustomFeatures()))])
```

```
# select and scale features
```

```
pipe_select_scale = Pipeline([("select", SelectKBest(score_func = chi2)),  
                               ("scale", StandardScaler(with_mean = False))])
```

```
[101]: # extract features
```

```
# you can also use bi-grams:
```

```
X_extract = pipe_extract.set_params(terms__tfidf__ngram_range = (1,2)).  
    ↪ fit_transform(X_train, y_train)
```

```
[102]: print(X_extract)
```

```
(0, 697)      0.07090144070985319  
(0, 745)      0.17483279959508186  
(0, 785)      0.043180618930796104  
(0, 814)      0.17483279959508186  
(0, 1163)     0.1648630344692586  
(0, 1895)     0.14402927962636455  
(0, 1900)     0.17483279959508186  
(0, 1940)     0.1648630344692586  
(0, 1942)     0.17483279959508186  
(0, 2001)     0.1523026158602164  
(0, 2004)     0.17483279959508186  
(0, 2318)     0.1648630344692586  
(0, 2320)     0.17483279959508186  
(0, 2605)     0.1307761823598249  
(0, 2609)     0.17483279959508186  
(0, 2799)     0.157789373540365  
(0, 2800)     0.17483279959508186  
(0, 3315)     0.13525918980549953  
(0, 3323)     0.17483279959508186  
(0, 3509)     0.09334204887662222  
(0, 3529)     0.12370252143093134  
(0, 4020)     0.17483279959508186  
(0, 4021)     0.17483279959508186  
(0, 4343)     0.1523026158602164  
(0, 4346)     0.17483279959508186  
:  
(898, 1293)   0.14945135091427067  
(898, 1311)   0.25345986443799984  
(898, 2151)   0.23900642479096382  
(898, 2153)   0.25345986443799984  
(898, 3706)   0.20880316378705438  
(898, 3708)   0.25345986443799984  
(898, 3882)   0.20880316378705438  
(898, 3884)   0.25345986443799984
```

(898, 3923)	0.1771738498926844
(898, 3935)	0.23900642479096382
(898, 4862)	0.1334273238341794
(898, 4894)	0.25345986443799984
(898, 6233)	0.09131789001952145
(898, 6325)	0.25345986443799984
(898, 6480)	0.23900642479096382
(898, 6481)	0.25345986443799984
(898, 6700)	0.0829510509697508
(898, 6721)	0.20404322788705004
(898, 6857)	0.25345986443799984
(898, 6858)	0.25345986443799984
(898, 6955)	0.18673654039843596
(898, 8149)	158.0
(898, 8151)	4.0
(899, 3095)	1.0
(899, 8149)	5.0

```
[103]: # extract all features
X_select_scale = pipe_select_scale.set_params(select__k = 500).
    ↪ fit_transform(X_extract, y_train)
print(X_select_scale)
```

(0, 222)	1.7463975570695727
(0, 224)	4.868401476553422
(0, 332)	3.3310922794931095
(0, 391)	3.8799422337719514
(0, 457)	4.664456373753558
(0, 498)	2.431404727966195
(0, 499)	0.6518039044082895
(1, 215)	1.955545214342361
(1, 216)	4.391664155129969
(1, 498)	0.7954037771785323
(1, 499)	0.32590195220414475
(2, 190)	3.9374599587618193
(2, 415)	2.9160268437218404
(2, 498)	2.205437745813203
(2, 499)	2.607215617633158
(3, 30)	3.467617207975343
(3, 35)	19.81764758098014
(3, 44)	4.029061322449084
(3, 254)	17.490606015887288
(3, 283)	21.01627223472825
(3, 294)	1.5146934286923344
(3, 295)	9.407246305477612
(3, 487)	4.614820388219279
(3, 498)	1.4371500064930298
(3, 499)	0.6518039044082895

```

:      :
(893, 494)    2.240618278439294
(893, 498)    1.6992717057905007
(893, 499)    0.6518039044082895
(894, 347)    2.7889851989069125
(894, 498)    0.831558494323011
(895, 211)    2.8248704708645893
(895, 397)    1.0910163308028165
(895, 498)    1.8529292536545352
(895, 499)    0.9777058566124343
(896, 72)     24.41599739893815
(896, 294)    1.6522845847666836
(896, 309)    16.430593476368912
(896, 352)    2.675593080916396
(896, 380)    4.5285814862791804
(896, 386)    13.6708637361617
(896, 498)    1.202144345053918
(897, 13)     2.5695129210768552
(897, 191)    6.145400961179911
(897, 498)    1.8438905743684155
(897, 499)    1.303607808816579
(898, 397)    1.4938347656108477
(898, 498)    1.4281113272069101
(898, 499)    1.303607808816579
(899, 193)    11.497860975131363
(899, 498)    0.04519339643059842

```

7.5.1 Using cross-validation with parameters (grid)

```

[104]: # Select best hyperparameters by cross validation
from sklearn.model_selection import GridSearchCV

# Model
logistic = LogisticRegression(max_iter=10000, tol=0.1, solver='lbfgs')

# Parameters: (np.logspace returns numbers spaced evenly on a log scale.)
param_logistic = {
    'C': np.logspace(-4, 4, 4)
}

# For an explanation of the 'C' parameter in scikit-learn logistic regression,
→ see:
# https://stackoverflow.com/questions/22851316/
→ what-is-the-inverse-of-regularization-strength-in-logistic-regression-how-shoul
# C= 1/\lambda where \lambda can be assimilated to the regularization parameter
# you probably have seen in the lasso regression
# Grid Search

```

```
cv_logistic = GridSearchCV(logistic, param_logistic, cv=10, scoring='f1')
cv_logistic.fit(X_select_scale, y_train)
```

```
[104]: GridSearchCV(cv=10, estimator=LogisticRegression(max_iter=10000, tol=0.1),
                param_grid={'C': array([1.00000000e-04, 4.64158883e-02,
                2.15443469e+01, 1.00000000e+04])},
                scoring='f1')
```

```
[105]: # See https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.
        ↪GridSearchCV.html
        print(cv_logistic.best_estimator_)
```

```
LogisticRegression(C=0.046415888336127774, max_iter=10000, tol=0.1)
```

```
[106]: print(cv_logistic.best_score_)
```

```
0.9020406477845386
```

7.5.2 Similar with pipeline

```
[107]: # Pipe Logistic
pipe_logistic = Pipeline([('select_scale', pipe_select_scale),
                           ('classify', LogisticRegression(max_iter=10000, tol=0.
                           ↪1, solver='lbfgs'))])

# Parameters of pipelines can be set using '__' separated parameter names:
param_logistic = {
    'classify__C': np.logspace(-4, 4, 3),
    'select_scale__select__k': [600, 1000, 5000]
}

cv_logistic = GridSearchCV(pipe_logistic, param_logistic, cv=10, scoring='f1')
cv_logistic.fit(X_extract, y_train)
```

```
[107]: GridSearchCV(cv=10,
                    estimator=Pipeline(steps=[('select_scale',
                                                Pipeline(steps=[('select',
                                                                    SelectKBest(k=500,
                                                                    score_func=<function chi2 at 0x7fe565f9b1f0>)),
                                                                    ('scale',
                                                                    StandardScaler(with_mean=False))])),
                                                ('classify',
                                                LogisticRegression(max_iter=10000,
                                                                    tol=0.1))]),
                    param_grid={'classify__C': array([1.e-04, 1.e+00, 1.e+04]),
                                'select_scale__select__k': [600, 1000, 5000]},
                    scoring='f1')
```



```
[108]: print(cv_logistic.best_estimator_)
```

```
Pipeline(steps=[('select_scale',
                  Pipeline(steps=[('select',
                                   SelectKBest(k=5000,
                                                score_func=<function chi2 at
0x7fe565f9b1f0>)),
                                  ('scale', StandardScaler(with_mean=False))])),
                ('classify',
                 LogisticRegression(C=10000.0, max_iter=10000, tol=0.1))])
```

```
[109]: print(cv_logistic.best_score_)
```

```
0.8457155973785053
```

7.6 Other Models

7.6.1 Naive Bayes

```
[110]: from sklearn.naive_bayes import MultinomialNB
```

```
# Pipe NB
pipe_nb = Pipeline([('select_scale', pipe_select_scale),
                    ('classify', MultinomialNB())])

# Parameters of pipelines can be set using '__' separated parameter names:
param_nb = {
    'classify__alpha': [0.5, 1, 10],
    'select_scale__select__k': [600, 1000, 5000]
}

cv_nb = GridSearchCV(pipe_nb, param_nb, cv=10, scoring='f1')
cv_nb.fit(X_extract, y_train)
print(cv_nb.best_score_)
```

```
0.8301298848126655
```

```
[111]: # full pipeline
model = Pipeline([("extract", FeatureUnion([("terms", Pipeline([('clean',
    ↳ CleanText()),
                                                                    ('tfidf',
    ↳ TfidfVectorizer(ngram_range = (1,2))])),
    ↳ ("custom", CustomFeatures())))),
                ("select", SelectKBest(score_func = chi2, k = 1000)),
                ("scale", StandardScaler(with_mean = False)),
                ("classify", MultinomialNB())])

# fitting
```

```

model.fit(X_train, y_train)

# final evaluation
y_pred = model.predict(X_test)
f1_score(y_test, y_pred)

```

[111]: 0.782608695652174

```

[112]: # we are now able to classify new messages!
#msg = pd.DataFrame(columns = ["text"],
#data      = ["REMINDER FROM O2: To get 2.50 pounds free call
↳credit and details of great offers pls reply 2 this text with your valid
↳name, house no and postcode"])

msg = pd.DataFrame(columns = ["text"],
data      = ["The bikes are heavy and unwieldy. The collection
↳and return of the bicycle is super-comfortable because the bikes are heavy"])

model.predict(msg)

```

[112]: array([0])

```

[113]: # we are now able to classify new messages!
#msg = pd.DataFrame(columns = ["text"],
#data      = ["REMINDER FROM O2: To get 2.50 pounds free call
↳credit and details of great offers pls reply 2 this text with your valid
↳name, house no and postcode"])

msg = pd.DataFrame(columns = ["text"],
data      = ["Satisfied"])

model.predict(msg)

```

[113]: array([1])

7.6.2 Support Vector Machine

```

[114]: from sklearn.svm import LinearSVC

# Pipe SVC
pipe_svc = Pipeline([('select_scale', pipe_select_scale),
                      ('classify', LinearSVC(max_iter=10000, tol=0.1))])

# Parameters of pipelines can be set using '__' separated parameter names:
param_svc = {
    'classify__C': [0.01, 0.1, 1],

```

```

        'select_scale__select__k': [600, 1000, 5000]
    }

    cv_svc = GridSearchCV(pipe_svc, param_svc, cv=10, scoring='f1')
    cv_svc.fit(X_extract, y_train)
    print(cv_svc.best_score_)

```

0.8463498495228174

```

[115]: # full pipeline
model = Pipeline([("extract", FeatureUnion([("terms", Pipeline([('clean',
    ↳ CleanText()),

                                                                    ('tfidf',
    ↳ TfidfVectorizer(ngram_range = (1,2)))))),

                                                                    ("custom", CustomFeatures())))),
                  ("select", SelectKBest(score_func = chi2, k = 1000)),
                  ("scale", StandardScaler(with_mean = False)),
                  ("classify", LinearSVC(C = 1, max_iter=10000, tol=0.1))])

# fitting
model.fit(X_train, y_train)

# final evaluation
y_pred = model.predict(X_test)
f1_score(y_test, y_pred)

```

[115]: 0.7972972972972973

```

[116]: # we are now able to classify new messages!
#msg = pd.DataFrame(columns = ["text"],
                        #data    = ["REMINDER FROM 02: To get 2.50 pounds free call
    ↳ credit and details of great offers pls reply 2 this text with your valid
    ↳ name, house no and postcode"])

msg = pd.DataFrame(columns = ["text"],
                    data    = ["The bikes are heavy and unwieldy. The collection
    ↳ and return of the bicycle is super-comfortable because the bikes are heavy"])

model.predict(msg)

```

[116]: array([0])

```

[117]: # we are now able to classify new messages!
#msg = pd.DataFrame(columns = ["text"],
                        #data    = ["REMINDER FROM 02: To get 2.50 pounds free call
    ↳ credit and details of great offers pls reply 2 this text with your valid
    ↳ name, house no and postcode"])

```

```
msg = pd.DataFrame(columns = ["text"],
                      data    = ["Satisfied"])

model.predict(msg)
```

[117]: array([1])

7.6.3 Random Forest

```
[118]: from sklearn.ensemble import RandomForestClassifier

# Pipe RF
pipe_rf = Pipeline([('select_scale', pipe_select_scale),
                    ('classify', RandomForestClassifier())])

# Parameters of pipelines can be set using '__' separated parameter names:
param_rf = {
    'classify__n_estimators': [100, 200],
    'select_scale__select__k': [600, 1000]
}

cv_rf = GridSearchCV(pipe_rf, param_rf, cv=10, scoring='f1')
cv_rf.fit(X_extract, y_train)
print(cv_rf.best_score_)
```

0.850498859773985

```
[119]: # full pipeline
model = Pipeline([("extract", FeatureUnion([("terms", Pipeline([('clean',
    ↳ CleanText()),
                                                                    ('tfidf',
    ↳ TfidfVectorizer(ngram_range = (1,2)))])),
                  ("custom", CustomFeatures()))]),
                  ("select", SelectKBest(score_func = chi2, k = 1000)),
                  ("scale", StandardScaler(with_mean = False)),
                  ("classify", RandomForestClassifier())])

# fitting
model.fit(X_train, y_train)

# final evaluation
y_pred = model.predict(X_test)
f1_score(y_test, y_pred)
```

[119]: 0.8079470198675497

```
[120]: # we are now able to classify new messages!
#msg = pd.DataFrame(columns = ["text"],
#data      = ["REMINDER FROM 02: To get 2.50 pounds free call_
↳credit and details of great offers pls reply 2 this text with your valid_
↳name, house no and postcode"])

msg = pd.DataFrame(columns = ["text"],
data      = ["The bikes are heavy and unwieldy. The collection_
↳and return of the bicycle is super-comfortable because the bikes are heavy"])

model.predict(msg)
```

```
[120]: array([1])
```

```
[121]: # we are now able to classify new messages!
#msg = pd.DataFrame(columns = ["text"],
#data      = ["REMINDER FROM 02: To get 2.50 pounds free call_
↳credit and details of great offers pls reply 2 this text with your valid_
↳name, house no and postcode"])

msg = pd.DataFrame(columns = ["text"],
data      = ["Satisfied"])

model.predict(msg)
```

```
[121]: array([1])
```

7.7 Long Example 1: Text clustering

```
[122]: import re
import string
import pandas as pd
```

```
[123]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.cluster import KMeans
from sklearn.cluster import AgglomerativeClustering
```

7.8 Full text for clustering

This corpus contain some strings about Google and some strings about TF-IDF from Wikipedia. Just for example

```
[124]: all_text = ""
Google and Facebook are strangling the free press to death. Democracy is the_
↳loser
Your 60-second guide to security stuff Google touted today at Next '18
```

```

A Guide to Using Android Without Selling Your Soul to Google
Review: Lenovo's Google Smart Display is pretty and intelligent
Google Maps user spots mysterious object submerged off the coast of Greece -
↳and no-one knows what it is
Android is better than IOS
In information retrieval, tf-idf or TFIDF, short for term frequency-inverse
↳document frequency
is a numerical statistic that is intended to reflect how important
a word is to a document in a collection or corpus.
It is often used as a weighting factor in searches of information retrieval
text mining, and user modeling. The tf-idf value increases proportionally
to the number of times a word appears in the document
and is offset by the frequency of the word in the corpus
"".split("\n")[1:-1]

```

```
[125]: all_text
```

```

[125]: ['Google and Facebook are strangling the free press to death. Democracy is the
loser',
'Your 60-second guide to security stuff Google touted today at Next '18",
'A Guide to Using Android Without Selling Your Soul to Google',
'Review: Lenovo's Google Smart Display is pretty and intelligent',
'Google Maps user spots mysterious object submerged off the coast of Greece -
and no-one knows what it is',
'Android is better than IOS',
'In information retrieval, tf-idf or TFIDF, short for term frequency-inverse
document frequency',
'is a numerical statistic that is intended to reflect how important ',
'a word is to a document in a collection or corpus.',
'It is often used as a weighting factor in searches of information retrieval',
'text mining, and user modeling. The tf-idf value increases proportionally',
'to the number of times a word appears in the document',
'and is offset by the frequency of the word in the corpus']

```

7.9 Preprocessing and tokenizing

Firstly, we must bring every chars to lowercase and remove all punctuation, because it's not important for our task, but is very harmful for clustering algorithm. After that, we'll split strings to array of words.

```

[126]: def preprocessing(line):
        line = line.lower()
        line = re.sub(r"[{}]" .format(string.punctuation), " ", line)
        return line

```

Now, let's calculate tf-idf for this corpus

```
[127]: tfidf_vectorizer = TfidfVectorizer(preprocessor=preprocessing)
tfidf = tfidf_vectorizer.fit_transform(all_text)
```

7.10 K-means

```
[128]: kmeans = KMeans(n_clusters=2)
```

```
[129]: list(zip(kmeans.fit_predict(tfidf), all_text))
```

```
[129]: [(1,
        'Google and Facebook are strangling the free press to death. Democracy is the
        loser'),
        (1, "Your 60-second guide to security stuff Google touted today at Next '18"),
        (1, 'A Guide to Using Android Without Selling Your Soul to Google'),
        (1, 'Review: Lenovo’s Google Smart Display is pretty and intelligent'),
        (1,
        'Google Maps user spots mysterious object submerged off the coast of Greece -
        and no-one knows what it is'),
        (1, 'Android is better than IOS'),
        (0,
        'In information retrieval, tf-idf or TFIDF, short for term frequency-inverse
        document frequency'),
        (1, 'is a numerical statistic that is intended to reflect how important '),
        (0, 'a word is to a document in a collection or corpus.'),
        (0,
        'It is often used as a weighting factor in searches of information
        retrieval'),
        (0,
        'text mining, and user modeling. The tf-idf value increases proportionally'),
        (0, 'to the number of times a word appears in the document'),
        (0, 'and is offset by the frequency of the word in the corpus')]
```

7.11 Agglomerative Clustering

```
[130]: hac = AgglomerativeClustering(n_clusters=2, affinity='cosine',
→linkage='average')
```

```
[131]: list(zip(hac.fit_predict(tfidf.toarray()), all_text))
```

```
[131]: [(0,
        'Google and Facebook are strangling the free press to death. Democracy is the
        loser'),
        (1, "Your 60-second guide to security stuff Google touted today at Next '18"),
        (1, 'A Guide to Using Android Without Selling Your Soul to Google'),
        (0, 'Review: Lenovo’s Google Smart Display is pretty and intelligent'),
        (0,
```

```

'Google Maps user spots mysterious object submerged off the coast of Greece -
and no-one knows what it is'),
(1, 'Android is better than IOS'),
(0,
'In information retrieval, tf-idf or TFIDF, short for term frequency-inverse
document frequency'),
(1, 'is a numerical statistic that is intended to reflect how important '),
(0, 'a word is to a document in a collection or corpus.'),
(0,
'It is often used as a weighting factor in searches of information
retrieval'),
(0,
'text mining, and user modeling. The tf-idf value increases proportionally'),
(0, 'to the number of times a word appears in the document'),
(0, 'and is offset by the frequency of the word in the corpus']]

```

7.12 Example 2: Topic model (1): BikeMi survey

```

[134]: import nltk
nltk.download('wordnet')

```

```

[nltk_data] Downloading package wordnet to
[nltk_data] /Users/giancarlomanzi/nltk_data...
[nltk_data] Package wordnet is already up-to-date!

```

```

[134]: True

```

7.12.1 Cleaning and pre-processing

```

[135]: from nltk.corpus import stopwords
from nltk.stem.wordnet import WordNetLemmatizer
import string
stop=set(stopwords.words('english'))
exclude=set(string.punctuation)
lemma=WordNetLemmatizer()
def clean(doc):
    stop_free=" ".join([i for i in doc.lower().split() if i not in stop])
    punc_free=''.join(ch for ch in stop_free if ch not in exclude)
    normalized=" ".join(lemma.lemmatize(word) for word in punc_free.split())
    return normalized

```

```

[136]: import pandas as pd
df = pd.read_csv('Polarity2014Reduced.csv', sep = ";", header = 0)
df.columns=['review','sentiment']
df2=df[df['sentiment']==-1]
df2.shape

```



```
[136]: (354, 2)
```

```
[137]: doc_complete=df2.iloc[0:2065,0].values.tolist()
doc_clean=[clean(doc).split() for doc in doc_complete]
```

8 Getting the document-term matrix

```
[138]: from sklearn.feature_extraction.text import CountVectorizer
import numpy as np
SOME_FIXED_SEED = 42
np.random.seed(SOME_FIXED_SEED)
```

```
[139]: cv=CountVectorizer(min_df=2,max_df=50,ngram_range=(1,2), token_pattern=None,
↳tokenizer=lambda doc:doc,preprocessor=lambda doc:doc)
```

```
[140]: cv_features=cv.fit_transform(doc_clean)
print(cv_features.shape)
vocabulary=np.array(cv.get_feature_names())
```

```
(354, 1392)
```

```
[141]: vocabulary
```

```
[141]: array(['1', '1 volta', '10', ..., '√® stato', '√® troppo', '√® un'],
dtype='<U24')
```

```
[142]: vocabulary
```

```
[142]: array(['1', '1 volta', '10', ..., '√® stato', '√® troppo', '√® un'],
dtype='<U24')
```

9 LDA ANALYSIS

```
[143]: # Using sklearn.decomposition LDA with 11 topics
from sklearn.decomposition import LatentDirichletAllocation
TOTAL_TOPICS=11
```

```
[144]: lda_model=LatentDirichletAllocation(n_components=TOTAL_TOPICS,max_iter=500,max_doc_update_iter
↳,random_state=42,n_jobs=16)
```

```
[145]: # Using the transformer 'fit_transform'
document_topics=lda_model.fit_transform(cv_features)
```

```
[146]: document_topics.shape
```

[146]: (354, 11)

```
[145]: # Extraqcting the most important 10 terms for each topic
topic_terms=lda_model.components_
top_terms=10 # number of 'top terms'
topic_key_terms_idx=np.argsort(-np.absolute(topic_terms), axis=1)[:,:top_terms]
topic_keyterms=vocabulary[topic_key_terms_idx]
topics=['', '.join(topic) for topic in topic_keyterms]
pd.set_option('display.max_colwidth',-1)
topics_df=pd.DataFrame(topics,columns=['Term per Topic'], index=['Topic'+str(t)
↳for t in range(1,TOTAL_TOPICS+1)])
topics_df
```

<ipython-input-145-515db2202b96>:7: FutureWarning: Passing a negative integer is deprecated in version 1.0 and will not be supported in future version. Instead, use None to not limit the column width.

```
pd.set_option('display.max_colwidth',-1)
```

[145]:

	Term per Topic
Topic1	della, completamente, servizio, la bici, possibilit√†, possibilit√† di, tramite, segnalare, app, troppo
Topic2	servizio, bicicletta, lasciare, la bicicletta, cambio, stalli, se, le stazioni, lasciare la, la bici
Topic3	piene, con, sempre, al, vuote, molto, tutte, alcune, le colonnine, colonnine
Topic4	mi, da, mi √®, servizio, stazione, la bici, se, √® capitato, capitato, bikemi
Topic5	con, tempo, migliorare, al, ecc, cambio, sistema, delle bici, migliorare il, anche
Topic6	essere, con, le bici, frequenza, essere pi√, manutenzione, bike, bici con, bici sono, troppo
Topic7	pi√ spesso, al, gomme, spesso le, della, cambio, le bici, del, gonfiare, controllare
Topic8	possibilit√† di, possibilit√†, segnalare, della, di segnalare, che non, dei, controllare pi√, delle biciclette, al
Topic9	troppo, piste, piste ciclabili, ciclabili, cambio, con, problemi, ruote, manutenzione, sgonfie
Topic10	anche, le stazioni, stazione, molto, piene, servizio, tempo, del, da, cambio
Topic11	stalli, gli, centro, gli stalli, ci, nelle, di punta, punta, aumentare, le bici

```
[146]: dt_df=pd.DataFrame(document_topics,columns=['T'+str(i) for i in
↳range(1,TOTAL_TOPICS+1)])
dt_df
```

```
[146]:
```

	T1	T2	T3	T4	T5	T6	T7 \
0	0.005348	0.005348	0.005348	0.005348	0.005348	0.005348	0.005348
1	0.009091	0.909084	0.009091	0.009091	0.009091	0.009091	0.009092
2	0.002841	0.971589	0.002841	0.002841	0.002841	0.002841	0.002841
3	0.015152	0.015153	0.015154	0.848456	0.015152	0.015160	0.015157
4	0.001567	0.001567	0.984325	0.001567	0.001567	0.001567	0.001567
..
349	0.003637	0.003637	0.003636	0.003637	0.003636	0.003637	0.003637
350	0.010101	0.010102	0.010102	0.898987	0.010101	0.010101	0.010102
351	0.018183	0.018182	0.018182	0.018182	0.018183	0.018182	0.018183
352	0.002392	0.002393	0.002393	0.002392	0.002392	0.002392	0.002392
353	0.011365	0.011364	0.011364	0.011364	0.886355	0.011365	0.011365

	T8	T9	T10	T11
0	0.005348	0.946523	0.005348	0.005348
1	0.009092	0.009093	0.009091	0.009091
2	0.002841	0.002841	0.002841	0.002841
3	0.015152	0.015156	0.015152	0.015157
4	0.001567	0.001568	0.001567	0.001568
..
349	0.003637	0.433846	0.003637	0.533425
350	0.010101	0.010101	0.010101	0.010101
351	0.818175	0.018182	0.018182	0.018182
352	0.002392	0.002393	0.976075	0.002392
353	0.011365	0.011364	0.011364	0.011364

[354 rows x 11 columns]

```
[147]: # Column 'Contribution%' gives the max probability among the 354
# features (terms) for each topic
dt_df=pd.DataFrame(document_topics,columns=['T'+str(i) for i in_
→range(1,TOTAL_TOPICS+1)])
pd.options.display.float_format='{:,.5f}'.format
pd.set_option('display.max_colwidth',200)
max_contrib_topics=dt_df.max(axis=0)
dominant_topics=max_contrib_topics.index
contrib_perc=max_contrib_topics.values
document_numbers=[dt_df[dt_df[t]==max_contrib_topics.loc[t]].index[0] for t in_
→dominant_topics]
results_df=pd.DataFrame({'Dominant Topic':dominant_topics,'Contribution%':
→contrib_perc, 'Answer Num': document_numbers,'Topic':topics_df['Term per_
→Topic']})
results_df
```

```
[147]:
```

	Dominant Topic	Contribution%	Answer Num \
Topic1	T1	0.97159	193
Topic2	T2	0.99209	179

Topic3	T3	0.98510	52
Topic4	T4	0.98978	328
Topic5	T5	0.99072	28
Topic6	T6	0.96503	126
Topic7	T7	0.98557	342
Topic8	T8	0.96503	174
Topic9	T9	0.98943	294
Topic10	T10	0.98864	198
Topic11	T11	0.99126	114

Topic

Topic1 della, completamente, servizio, la bici, possibilit√†, possibilit√† di, tramite, segnalare, app, troppo

Topic2 servizio, bicicletta, lasciare, la bicicletta, cambio, stalli, se, le stazioni, lasciare la, la bici

Topic3 piene, con, sempre, al, vuote, molto, tutte, alcune, le colonnine, colonnine

Topic4 mi, da, mi √®, servizio, stazione, la bici, se, √® capitato, capitato, bikemi

Topic5 con, tempo, migliorare, al, ecc, cambio, sistema, delle bici, migliorare il, anche

Topic6 essere, con, le bici, frequenza, essere pi√, manutenzione, bike, bici con, bici sono, troppo

Topic7 pi√ spesso, al, gomme, spesso le, della, cambio, le bici, del, gonfiare, controllare

Topic8 possibilit√† di, possibilit√†, segnalare, della, di segnalare, che non, dei, controllare pi√, delle biciclette, al

Topic9 troppo, piste, piste ciclabili, ciclabili, cambio, con, problemi, ruote, manutenzione, sgonfie

Topic10 anche, le stazioni, stazione, molto, piene, servizio, tempo, del, da, cambio

Topic11 stalli, gli, centro, gli stalli, ci, nelle, di punta, punta, aumentare, le bici

```
[148]: # This gives, for each topic, the % of features having prob >0.9
numT1=np.count_nonzero(dt_df['T1']>0.9)
FrT1=numT1/2133
numT2=np.count_nonzero(dt_df['T2']>0.9)
FrT2=numT2/2133
numT3=np.count_nonzero(dt_df['T3']>0.9)
FrT3=numT3/2133
numT4=np.count_nonzero(dt_df['T4']>0.9)
FrT4=numT4/2133
numT5=np.count_nonzero(dt_df['T5']>0.9)
FrT5=numT5/2133
numT6=np.count_nonzero(dt_df['T6']>0.9)
FrT6=numT6/2133
```

```

numT7=np.count_nonzero(dt_df['T7']>0.9)
FrT7=numT7/2133
numT8=np.count_nonzero(dt_df['T8']>0.9)
FrT8=numT8/2133
numT9=np.count_nonzero(dt_df['T9']>0.9)
FrT9=numT9/2133
numT10=np.count_nonzero(dt_df['T10']>0.9)
FrT10=numT10/2133
numT11=np.count_nonzero(dt_df['T11']>0.9)
FrT11=numT11/2133
d=(FrT1,FrT2,FrT3,FrT4,FrT5,FrT6,FrT7,FrT8,FrT9,FrT10,FrT11)
df_Fr=pd.DataFrame(data=d)
results_df.insert(4,'Freq 0.9-1',df_Fr.values)
results_df

```

```

[148]:
Dominant Topic  Contribution%  Answer Num  \
Topic1          T1           0.97159         193
Topic2          T2           0.99209         179
Topic3          T3           0.98510          52
Topic4          T4           0.98978        328
Topic5          T5           0.99072          28
Topic6          T6           0.96503        126
Topic7          T7           0.98557        342
Topic8          T8           0.96503        174
Topic9          T9           0.98943        294
Topic10         T10           0.98864        198
Topic11         T11           0.99126        114

```

```

Topic  \
Topic1      della, completamente, servizio, la bici, possibilit√,
possibilit√ di, tramite, segnalare, app, troppo
Topic2      servizio, bicicletta, lasciare, la bicicletta, cambio,
stalli, se, le stazioni, lasciare la, la bici
Topic3      piene, con, sempre, al, vuote,
molto, tutte, alcune, le colonnine, colonnine
Topic4      mi, da, mi √®, servizio,
stazione, la bici, se, √® capitato, capitato, bikemi
Topic5      con, tempo, migliorare, al, ecc,
cambio, sistema, delle bici, migliorare il, anche
Topic6      essere, con, le bici, frequenza, essere pi√,
manutenzione, bike, bici con, bici sono, troppo
Topic7      pi√ spesso, al, gomme, spesso le, della,
cambio, le bici, del, gonfiare, controllare
Topic8      possibilit√ di, possibilit√, segnalare, della, di segnalare, che non,
dei, controllare pi√, delle biciclette, al
Topic9      troppo, piste, piste ciclabili, ciclabili, cambio,
con, problemi, ruote, manutenzione, sgonfie

```

Topic10		anche, le stazioni, stazione,
molto, piene, servizio, tempo, del, da, cambio		
Topic11		stalli, gli, centro, gli stalli,
ci, nelle, di punta, punta, aumentare, le bici		

	Freq
Topic1	0.00516
Topic2	0.01547
Topic3	0.01172
Topic4	0.00891
Topic5	0.00609
Topic6	0.00469
Topic7	0.00797
Topic8	0.00797
Topic9	0.02391
Topic10	0.00656
Topic11	0.01828

```
[1]: #This is to let you have larger fonts...
from IPython.core.display import HTML
HTML("""
<style>

div.cell { /* Tunes the space between cells */
margin-top:1em;
margin-bottom:1em;
}

div.text_cell_render h1 { /* Main titles bigger, centered */
font-size: 2.2em;
line-height:1.4em;
text-align:center;
}

div.text_cell_render h2 { /* Parts names nearer from text */
margin-bottom: -0.4em;
}

div.text_cell_render { /* Customize text cells */
font-family: 'Times New Roman';
font-size:1.5em;
line-height:1.4em;
padding-left:3em;
padding-right:3em;
}

</style>
```

```
"""
```

```
[1]: <IPython.core.display.HTML object>
```

```
[ ]:
```